

# Quantitative Sports Betting Analysis System

Elo-Based Prediction, Monte Carlo Simulation,  
Expected Value Analysis, and Strategy Backtesting  
in the U.S. Sports Betting Market

Godfred Antwi Koduah

April 2026

Personal project in applied quantitative finance.  
Full Python source code and reproducible analysis pipeline available on request.

# Contents

|          |                                                               |           |
|----------|---------------------------------------------------------------|-----------|
| <b>1</b> | <b>Executive Summary</b>                                      | <b>3</b>  |
| <b>2</b> | <b>Statement of the Problem</b>                               | <b>4</b>  |
| 2.1      | Background and Motivation . . . . .                           | 4         |
| 2.2      | Specific Sub-Problems . . . . .                               | 4         |
| 2.3      | Scope and Limitations . . . . .                               | 5         |
| <b>3</b> | <b>Description of the Mathematics</b>                         | <b>6</b>  |
| 3.1      | Odds Conversion and Implied Probability . . . . .             | 6         |
| 3.1.1    | American Odds . . . . .                                       | 6         |
| 3.1.2    | Decimal Odds . . . . .                                        | 6         |
| 3.1.3    | Vigorish (Overround) . . . . .                                | 6         |
| 3.2      | Expected Value . . . . .                                      | 7         |
| 3.3      | Elo Rating System . . . . .                                   | 7         |
| 3.3.1    | Between-Season Regression . . . . .                           | 8         |
| 3.4      | Monte Carlo Simulation . . . . .                              | 9         |
| 3.5      | Parlay Compound Probability . . . . .                         | 10        |
| 3.6      | Kelly Criterion . . . . .                                     | 11        |
| 3.7      | Model Calibration Metrics . . . . .                           | 11        |
| <b>4</b> | <b>Description of the Algorithm</b>                           | <b>13</b> |
| 4.1      | System Architecture . . . . .                                 | 13        |
| 4.2      | Elo Update Procedure . . . . .                                | 13        |
| 4.3      | Synthetic Data Generation . . . . .                           | 14        |
| 4.4      | Backtesting Engine . . . . .                                  | 15        |
| 4.5      | Complexity Analysis . . . . .                                 | 16        |
| <b>5</b> | <b>Description of the Testing</b>                             | <b>17</b> |
| 5.1      | Unit Validation . . . . .                                     | 17        |
| 5.2      | Synthetic Data Validation . . . . .                           | 17        |
| 5.3      | Calibration Testing . . . . .                                 | 17        |
| 5.4      | Backtest Methodology . . . . .                                | 18        |
| <b>6</b> | <b>Results and Discussion</b>                                 | <b>19</b> |
| 6.1      | Parlay Risk Analysis . . . . .                                | 19        |
| 6.2      | Elo Rating Distribution . . . . .                             | 19        |
| 6.3      | Monte Carlo Season Projections . . . . .                      | 20        |
| 6.4      | Backtest Performance . . . . .                                | 21        |
| 6.4.1    | Interpretation: Why the Strategy Fails . . . . .              | 23        |
| 6.4.2    | What Would Be Needed for Profitability . . . . .              | 24        |
| 6.5      | Model Calibration . . . . .                                   | 24        |
| 6.6      | Kelly Criterion Bankroll Paths . . . . .                      | 25        |
| 6.7      | Vig Distribution and Market Structure . . . . .               | 26        |
| 6.8      | Sensitivity Analysis: Effect of Strategy Parameters . . . . . | 27        |
| 6.8.1    | Edge Threshold Sensitivity . . . . .                          | 27        |
| 6.8.2    | Kelly Fraction Sensitivity . . . . .                          | 27        |

|          |                                                                        |           |
|----------|------------------------------------------------------------------------|-----------|
| 6.8.3    | Elo K-Factor Sensitivity . . . . .                                     | 27        |
| 6.9      | Comparison to Financial Market Analogues . . . . .                     | 27        |
| 6.9.1    | Market Efficiency . . . . .                                            | 28        |
| 6.9.2    | The Vig as Bid-Ask Spread . . . . .                                    | 28        |
| 6.9.3    | Kelly Criterion in Portfolio Management . . . . .                      | 28        |
| 6.10     | Detailed Discussion of Elo Model Properties . . . . .                  | 28        |
| 6.10.1   | The Logistic Model as a Choice Function . . . . .                      | 28        |
| 6.10.2   | Convergence Properties . . . . .                                       | 29        |
| 6.10.3   | Home-Court Advantage Model . . . . .                                   | 29        |
| 6.11     | Limitations and Sources of Error . . . . .                             | 30        |
| 6.12     | Broader Implications . . . . .                                         | 30        |
| 6.13     | Extended Analysis: Poisson Scoring Model . . . . .                     | 31        |
| 6.13.1   | The Poisson Model for Game Scores . . . . .                            | 31        |
| 6.13.2   | Application to Point Spread Markets . . . . .                          | 32        |
| 6.14     | Closing Line Value (CLV) Analysis . . . . .                            | 32        |
| 6.15     | Information Edge Decay and Speed-to-Market . . . . .                   | 33        |
| 6.16     | Strategy Parameter Sensitivity Analysis . . . . .                      | 35        |
| 6.17     | Risk of Ruin Analysis . . . . .                                        | 36        |
| 6.18     | Cross-Asset Risk-Adjusted Return Comparison . . . . .                  | 37        |
| 6.19     | Arbitrage Detection Framework . . . . .                                | 38        |
| 6.20     | Practical Applications: Production System Architecture . . . . .       | 38        |
| 6.20.1   | Research Function . . . . .                                            | 38        |
| 6.20.2   | Trading Function . . . . .                                             | 39        |
| 6.20.3   | Technology Function . . . . .                                          | 39        |
| 6.21     | Market Microstructure: Price Discovery in Sports Betting . . . . .     | 41        |
| 6.21.1   | The Price Discovery Process . . . . .                                  | 41        |
| 6.21.2   | Closing Line as the Efficient Price . . . . .                          | 42        |
| 6.22     | Bayesian In-Game Updating: Live Betting Models . . . . .               | 42        |
| 6.22.1   | Mathematical Framework . . . . .                                       | 42        |
| 6.22.2   | Trading Implications . . . . .                                         | 43        |
| 6.23     | Portfolio Theory Applied to Sports Betting . . . . .                   | 44        |
| 6.23.1   | Correlation Structure of Game Outcomes . . . . .                       | 44        |
| 6.23.2   | Optimal Bet Allocation . . . . .                                       | 45        |
| 6.24     | Signal Analysis: Feature Importance and Half-Life Estimation . . . . . | 45        |
| 6.24.1   | Signal Half-Life . . . . .                                             | 45        |
| 6.25     | Market Maker Simulation: The Other Side of the Trade . . . . .         | 46        |
| 6.25.1   | Market Maker P&L Model . . . . .                                       | 47        |
| 6.25.2   | Implications for Line-Setting . . . . .                                | 47        |
| <b>7</b> | <b>Conclusion</b>                                                      | <b>49</b> |
| <b>A</b> | <b>Key Implementation Code Excerpts</b>                                | <b>53</b> |
| A.1      | Odds Converter . . . . .                                               | 53        |
| A.2      | Elo Rating System . . . . .                                            | 53        |
| A.3      | Kelly Criterion . . . . .                                              | 54        |
| A.4      | Monte Carlo Simulator . . . . .                                        | 54        |
| A.5      | Expected Value Calculator . . . . .                                    | 55        |
| A.6      | Model Calibration . . . . .                                            | 55        |

# 1 Executive Summary

This report presents a comprehensive quantitative analysis system for the U.S. sports betting market, developed to explore whether the quantitative tools of financial mathematics—probability theory, stochastic simulation, and risk management—can be systematically applied to the U.S. sports betting market. The project implements, tests, and evaluates a suite of mathematical models, translating theoretical concepts into working Python code with rigorous backtesting and empirical validation.

The system comprises ten interconnected modules: (1) an odds conversion and implied probability engine supporting American, decimal, and fractional odds formats with vigorish decomposition; (2) an expected value calculator for identifying positive-EV betting opportunities; (3) an Elo rating system for NBA team strength estimation with home-court advantage, margin-of-victory adjustment, and between-season regression; (4) a Monte Carlo season simulator using 10,000 iterations to generate playoff probability distributions; (5) a parlay risk analyzer demonstrating compound probability decay and super-linear house edge growth; (6) a market efficiency analyzer with calibration testing, Brier score, and log loss computation; (7) a Kelly Criterion position-sizing framework comparing full, half, quarter, and tenth Kelly fractions; (8) a synthetic data generator producing realistic NBA season data with controllable noise and vigorish levels; (9) a full backtesting engine implementing a walk-forward +EV strategy; and (10) a visualization engine producing seven multi-panel figures for the report.

The analysis processed 720 simulated games across 16 NBA teams. The Elo model produced well-calibrated win probabilities with a Brier score of 0.2312 (beating the trivial baseline of 0.2500) and log loss of 0.6550 (beating the baseline of 0.6931). Monte Carlo simulation of 10,000 seasons generated robust playoff probability distributions with standard errors below 0.5%. The parlay analysis confirmed that the effective house edge compounds from 4.8% on a single bet to 26.4% on a 5-leg parlay and 63.2% on a 10-leg parlay, providing rigorous quantitative confirmation of industry claims about parlay profitability.

The +EV backtesting strategy, using quarter-Kelly position sizing with a 3% edge threshold, placed 531 bets with a win rate of 48.0% and an ROI of  $-5.81\%$ . The negative backtest result is itself a critically important finding: it demonstrates that a basic Elo model, while well-calibrated for probability estimation, does not generate sufficient edge to overcome the sportsbook's vigorish in a realistic market environment. This is the sports-betting analogue of the efficient market hypothesis—closing lines incorporate substantial information, and only models with genuine informational advantages can sustainably extract positive returns. The report discusses in detail the implications of this result and the specific model enhancements that would be required to achieve profitability.

All computations are performed in IEEE 754 double-precision floating-point arithmetic. The implementation uses only NumPy and Matplotlib as external dependencies, with all mathematical models coded from first principles. Seven figures with 17 total panels are produced, each discussed in detail in the Results section.

## 2 Statement of the Problem

### 2.1 Background and Motivation

The U.S. sports betting industry has experienced explosive growth since the Supreme Court struck down the Professional and Amateur Sports Protection Act (PASPA) in *Murphy v. National Collegiate Athletic Association*, 584 U.S. 453 (2018). The industry has processed over \$625 billion in cumulative handle (total amount wagered), generated over \$62 billion in total operator revenue, and the global online gambling market exceeds \$80 billion. The market is projected to reach \$150 billion by 2030, growing at a compound annual growth rate of approximately 10%.

There exists a fundamental structural analogy between sports betting markets and financial markets. Sportsbooks function as market makers, embedding a margin (vigorous, or “vig”) into their published odds that guarantees profit regardless of outcome—structurally identical to the bid-ask spread charged by a broker-dealer in equity markets. The “alpha” sought by quantitative traders in financial markets corresponds to “positive expected value” (positive EV) in sports betting: both represent systematic edges over the market’s implied probabilities.

This analogy raises a natural quantitative question that forms the basis of this project: **can a mathematical model identify and exploit systematic mispricings in the sports betting market, and if not, what does that tell us about market efficiency?**

### 2.2 Specific Sub-Problems

More precisely, this project addresses the following five sub-problems, each requiring distinct mathematical tools and computational approaches:

**(a) Probability estimation.** Given historical game results, can an Elo-based rating system produce well-calibrated win probabilities for NBA games? The Elo system, originally developed by Arpad Elo for chess [2], assigns numerical ratings to teams that reflect relative strength. The key mathematical question is whether the logistic model underlying Elo—which maps rating differences to win probabilities—produces probability estimates that are *calibrated*: games predicted to be won 60% of the time should actually be won approximately 60% of the time. The reason calibration is the right metric (rather than, say, classification accuracy) is that profitable betting requires accurate probability estimates, not just correct binary predictions. A model that predicts the favorite will win every game may have high accuracy (favorites win ~60% of the time) but cannot identify value bets because it provides no information about the *magnitude* of the edge.

**(b) Market efficiency.** Do sportsbook closing lines incorporate enough information that a basic Elo model cannot consistently identify positive expected value opportunities? The reason this question matters is that the existence of a well-calibrated model is necessary but not sufficient for profitable betting: the model must also be *more* accurate than the market’s implied probabilities. If the market has already incorporated the information captured by Elo ratings (team win-loss records, home-court advantage, margin of victory), then the model’s probability estimates will differ from the market’s only by noise, not by systematic edge. The measure of market efficiency we use is the average vig (overround) across games, computed as the sum of implied probabilities minus one.

**(c) Strategy profitability.** If we implement a disciplined +EV betting strategy using Kelly Criterion position sizing, can we generate positive long-term returns against realistic market odds? The reason for using Kelly sizing (rather than flat betting) is that the Kelly Criterion maximizes the long-run geometric growth rate of the bankroll, which is the appropriate objective for a bettor with a fixed bankroll who reinvests profits. The strategy involves three parameters: the Elo model’s K-

factor (which controls how quickly ratings adapt to new results), the edge threshold (the minimum model-vs-market probability difference required to place a bet), and the Kelly fraction (which controls the aggressiveness of position sizing).

**(d) Risk quantification.** What is the variance, maximum drawdown, and risk-adjusted return profile of the +EV strategy? The reason risk quantification is important (beyond just measuring return) is that a strategy with positive expected value can still experience long losing streaks and significant drawdowns due to the inherent randomness of individual game outcomes. The maximum drawdown—the largest peak-to-trough decline in bankroll—is the most relevant risk metric for a bettor because it determines whether the strategy can survive periods of adverse variance without exhausting the bankroll. The Sharpe-like ratio (mean return divided by standard deviation of returns, scaled by the square root of the number of bets) provides a risk-adjusted measure analogous to the Sharpe ratio used in portfolio management.

**(e) Compound risk in parlays.** How does the house edge compound as the number of parlay legs increases, and what mathematical framework explains why sportsbooks earn their highest margins on multi-leg bets? The reason this analysis is important is that parlays represent a significant and growing share of sportsbook revenue, and the mathematical mechanism—multiplicative compounding of per-leg vig—is not intuitively obvious to most bettors. The analysis provides rigorous quantitative support for the widely-held claim that “the house loves parlays” and connects the compound probability framework to concepts in financial mathematics (e.g., the decay of joint probabilities in portfolio insurance).

## 2.3 Scope and Limitations

The project uses synthetic data generated from known team strength distributions rather than scraped historical data. The reason for this design choice is threefold. First, synthetic data provides full control over the data-generating process, allowing us to verify that the system behaves correctly when the true probabilities are known. Second, it ensures reproducibility: the random seed (2026) produces identical results on any machine with the same NumPy version. Third, it avoids the legal and ethical complications of web scraping from commercial sportsbook websites.

The primary limitation of synthetic data is that the market odds are generated by adding noise to the true probabilities rather than reflecting the complex information aggregation process of real sportsbook lines. Real closing lines incorporate public betting patterns, sharp money, injury reports, weather conditions, and proprietary models—information sources that our synthetic market does not capture. The consequence is that our backtest results may overstate or understate the difficulty of beating real markets, depending on whether the synthetic vig distribution matches the empirical distribution.

All computations are performed in IEEE 754 double-precision (64-bit) floating-point arithmetic, providing approximately 15–16 significant decimal digits and a machine epsilon of  $\epsilon_{\text{mach}} \approx 2.22 \times 10^{-16}$ . This precision level is more than sufficient for the financial calculations in this project, where the relevant quantities (probabilities, odds, bankroll values) typically have at most 4–6 significant figures.

### 3 Description of the Mathematics

This section presents the mathematical foundations underlying each module of the analysis system. The models draw from probability theory, statistical estimation, stochastic simulation, information theory, and financial mathematics. Each subsection develops the theory in full detail, explains the reason each formula takes its specific form, and connects the mathematics to the broader sports betting and financial markets context.

#### 3.1 Odds Conversion and Implied Probability

##### 3.1.1 American Odds

American odds are the standard format used by U.S. sportsbooks. They encode the payout structure of a bet relative to a \$100 baseline. A negative number (e.g.,  $-150$ ) indicates the amount a bettor must risk to win \$100; a positive number (e.g.,  $+130$ ) indicates the profit on a \$100 wager.

**Definition 3.1** (Implied Probability from American Odds). *For American odds  $o$ , the raw implied probability (which includes the vigorish) is:*

$$p_{\text{implied}} = \begin{cases} \frac{|o|}{|o| + 100} & \text{if } o < 0 \text{ (favorite),} \\ \frac{100}{o + 100} & \text{if } o > 0 \text{ (underdog).} \end{cases} \quad (1)$$

The reason the formulas differ for favorites and underdogs is rooted in the payout structure. For a favorite at  $-150$ : the bettor risks \$150 to win \$100, so the total return on a win is \$250 (stake + profit). The implied probability is  $150/250 = 0.60$ , which matches  $| -150 | / (| -150 | + 100) = 150/250$ . For an underdog at  $+130$ : the bettor risks \$100 to win \$130, total return \$230. The implied probability is  $100/230 = 0.4348$ , matching  $100/(130 + 100)$ .

##### 3.1.2 Decimal Odds

The conversion to decimal odds  $D$  is:

$$D = \begin{cases} 1 + \frac{100}{|o|} & \text{if } o < 0, \\ 1 + \frac{o}{100} & \text{if } o > 0. \end{cases} \quad (2)$$

The reason decimal odds are useful is that the payout on a winning bet is simply  $D \times \text{stake}$ , making expected value calculations more straightforward than with American odds.

##### 3.1.3 Vigorish (Overround)

In a two-outcome market (e.g., Team A vs. Team B), the sportsbook publishes odds for both sides. The vigorish is the excess of the implied probabilities over 1:

$$\text{vig} = p_A^{\text{implied}} + p_B^{\text{implied}} - 1. \quad (3)$$

The reason the vig is always positive in practice is that the sportsbook adjusts both sides' odds to ensure a guaranteed profit regardless of the outcome. The standard U.S. market is  $-110/-110$ , which gives  $p_A = 0.5238$  and  $p_B = 0.5238$ , so  $\text{vig} = 1.0476 - 1 = 0.0476$  or 4.76%. This 4.76%

overround is the sportsbook’s margin—analogous to the bid-ask spread in financial markets, as discussed in the market microstructure literature.

To recover *fair* (vig-free) probabilities from the market odds, we normalize:

$$p_A^{\text{fair}} = \frac{p_A^{\text{implied}}}{p_A^{\text{implied}} + p_B^{\text{implied}}}, \quad p_B^{\text{fair}} = \frac{p_B^{\text{implied}}}{p_A^{\text{implied}} + p_B^{\text{implied}}}. \quad (4)$$

The reason this normalization is the correct procedure (rather than, say, subtracting vig/2 from each side) is that it preserves the *ratio* of the two probabilities, which reflects the market’s assessment of relative team strength. The vig-removal method that subtracts equally from both sides would distort the ratio and could produce negative probabilities for extreme underdogs.

### 3.2 Expected Value

**Definition 3.2** (Expected Value of a Bet). *The expected value of a bet with true win probability  $p$ , decimal odds  $D$ , and stake  $S$  is:*

$$\text{EV} = p \cdot (D - 1) \cdot S - (1 - p) \cdot S = S[p \cdot D - 1]. \quad (5)$$

The reason the formula takes this form is straightforward from the definition of expected value: with probability  $p$  the bettor wins  $(D - 1) \cdot S$  (net profit), and with probability  $(1 - p)$  the bettor loses  $S$  (the stake). The simplified form  $S[pD - 1]$  makes it clear that a bet has positive expected value if and only if  $pD > 1$ , i.e., the true probability exceeds the reciprocal of the decimal odds.

**Definition 3.3** (Edge). *The edge of a model-based bet is the difference between the model’s estimated probability and the market’s implied probability:*

$$\text{edge} = p_{\text{model}} - p_{\text{implied}}. \quad (6)$$

The reason edge is defined relative to the implied probability (rather than the fair probability) is that the bettor faces the actual market odds including vig, not the fair odds. A bet with  $p_{\text{model}} > p_{\text{implied}}$  has positive expected value at the offered odds, even if the true probability is less than  $p_{\text{model}}$ , as long as the estimation error is not too large.

**Definition 3.4** (Breakeven Probability). *The breakeven probability at American odds  $o$  is the minimum win rate required to avoid losing money in the long run:*

$$p_{\text{break}} = p_{\text{implied}}(o) = \begin{cases} |o|/(|o| + 100) & o < 0, \\ 100/(o + 100) & o > 0. \end{cases} \quad (7)$$

The reason the breakeven probability equals the implied probability is that the implied probability is precisely the win rate at which the expected value is zero. At the standard  $-110$  odds, the breakeven is 52.38%: a bettor must win more than 52.38% of their bets just to break even, because they risk \$110 to win \$100 on each bet.

### 3.3 Elo Rating System

The Elo rating system, originally developed for chess by Arpad Elo [2] and later adapted for team sports by FiveThirtyEight [4], assigns each team a numerical rating  $R$  that reflects relative strength. The mathematical framework is based on the logistic distribution.



**Definition 3.5** (Expected Score). *The expected score (win probability) for team A against team B is:*

$$E(A) = \frac{1}{1 + 10^{(R_B - R_A^*)/400}}, \quad (8)$$

where  $R_A^* = R_A + H$  if team A is playing at home (with home-court advantage  $H$ ), and  $R_A^* = R_A$  otherwise.

The reason the formula uses base 10 and a scaling factor of 400 (rather than, say, the natural exponential  $e$  and a different scale) is historical convention from chess. The choice of base is mathematically irrelevant—changing the base simply rescales the ratings by a constant factor—but the convention  $10^{\Delta R/400}$  has the convenient property that a 400-point rating advantage corresponds to a 10:1 expected score ratio, i.e., a win probability of approximately 0.909. The logistic function  $1/(1 + 10^{-x})$  is a sigmoid curve that maps any rating difference  $\Delta R \in (-\infty, +\infty)$  to a probability in  $(0, 1)$ , ensuring that the output is always a valid probability regardless of the input.

**Definition 3.6** (Rating Update). *After a game between teams A and B with actual outcome  $S_A \in \{0, 1\}$  (1 for a win, 0 for a loss), the ratings update according to:*

$$R'_A = R_A + K_{adj} \cdot (S_A - E(A)), \quad R'_B = R_B + K_{adj} \cdot (S_B - E(B)), \quad (9)$$

where  $S_B = 1 - S_A$ ,  $E(B) = 1 - E(A)$ , and  $K_{adj}$  is the adjusted sensitivity parameter:

$$K_{adj} = K \cdot \ln(MOV + 1), \quad (10)$$

with  $K = 20$  being the base sensitivity and  $MOV = |\text{score}_A - \text{score}_B|$  the margin of victory.

The reason for the margin-of-victory adjustment (10) is that a team winning by 30 points provides stronger evidence of superiority than a team winning by 1 point. Without this adjustment, both outcomes would produce identical rating changes, wasting information. The logarithmic transformation  $\ln(MOV + 1)$  provides diminishing returns: the difference between a 1-point and a 10-point margin is weighted more heavily than the difference between a 30-point and a 39-point margin, reflecting the empirical observation that very large margins carry decreasing additional information about team quality. The reason for using  $\ln$  rather than a linear function  $K \cdot MOV$  is that linear scaling would allow a single blowout game to produce an unreasonably large rating change, destabilizing the system.

The Elo update has two critical mathematical properties. First, it is **zero-sum**: the total rating change across both teams is zero, since  $(S_A - E(A)) + (S_B - E(B)) = (S_A + S_B) - (E(A) + E(B)) = 1 - 1 = 0$ . This ensures that the average rating across all teams remains constant over time, providing a stable reference point. Second, the update is **self-correcting**: if a team is overrated (i.e.,  $E(A)$  is too high relative to  $S_A$ ), the update  $R'_A = R_A + K(S_A - E(A))$  decreases  $R_A$ , pulling it back toward its true level. This mean-reversion property makes the Elo system robust to temporary fluctuations.

### 3.3.1 Between-Season Regression

Between seasons, ratings regress toward the league mean:

$$R_{\text{new}} = \alpha \cdot R_{\text{old}} + (1 - \alpha) \cdot R_0, \quad (11)$$

where  $\alpha = 0.75$  is the reversion parameter and  $R_0 = 1500$  is the league average. The reason for this regression is that approximately 25% of a team's performance is not persistent across seasons due to

roster turnover, coaching changes, player development, and other factors. Setting  $\alpha = 0.75$  means that a team with rating 1600 at end-of-season starts the new season at  $0.75(1600) + 0.25(1500) = 1575$ , reflecting the empirical regularity that strong teams tend to regress somewhat while weak teams tend to improve.

### 3.4 Monte Carlo Simulation

Monte Carlo simulation estimates the distribution of season outcomes by repeatedly sampling from the stochastic process governing individual games. The mathematical framework is as follows.

Let  $\mathcal{S} = \{(h_k, a_k)\}_{k=1}^G$  denote the season schedule, where  $h_k$  and  $a_k$  are the home and away teams in game  $k$ , and  $G$  is the total number of games. For each simulation run  $s = 1, \dots, N$ :

1. For each game  $k$ , compute the win probability  $p_k = E(h_k)$  using the Elo model (8).
2. Draw a uniform random variable  $u_k \sim U(0, 1)$ . The home team wins if  $u_k < p_k$ ; otherwise the away team wins.
3. Accumulate win totals:  $W_{s,t} = \sum_{k=1}^G \mathbf{1}[\text{team } t \text{ wins game } k]$  for each team  $t$ .

After  $N$  simulations, the playoff probability for team  $t$  with threshold  $\tau$  wins is:

$$\hat{P}(\text{playoff} \mid t) = \frac{1}{N} \sum_{s=1}^N \mathbf{1}[W_{s,t} \geq \tau]. \quad (12)$$

**Proposition 3.1** (Monte Carlo Standard Error). *The standard error of the playoff probability estimate (12) is:*

$$SE = \sqrt{\frac{\hat{p}(1 - \hat{p})}{N}} \leq \frac{1}{2\sqrt{N}}, \quad (13)$$

where the upper bound is achieved when  $\hat{p} = 0.5$ . For  $N = 10,000$ , the maximum standard error is  $1/(2\sqrt{10000}) = 0.005 = 0.5\%$ .

The reason for using  $N = 10,000$  (rather than, say, 1,000 or 100,000) is a balance between precision and computational cost. At  $N = 1,000$ , the standard error would be 1.58%, meaning a team with true playoff probability 50% could report anywhere from 47% to 53% with one standard deviation of noise. At  $N = 10,000$ , this range narrows to 49.5%–50.5%, which is sufficient for practical use. At  $N = 100,000$ , the improvement (0.16% SE) would be negligible for decision-making while increasing computation time tenfold.

The convergence of the Monte Carlo estimator to the true probability is guaranteed by the Strong Law of Large Numbers: as  $N \rightarrow \infty$ ,  $\hat{P} \rightarrow P$  almost surely. The rate of convergence is  $O(N^{-1/2})$ —the standard Monte Carlo rate—regardless of the dimension of the problem (number of games, number of teams). This dimension-independence is the fundamental advantage of Monte Carlo over deterministic methods: for a season with  $G = 720$  games and 16 teams, the “dimension” of the probability space is 720 (one binary outcome per game), making deterministic enumeration of all  $2^{720}$  outcomes computationally infeasible.

The simulation uses *frozen* Elo ratings (not updating during simulation) as the input probability model. The reason for freezing is that the purpose of the simulation is to project future outcomes given current team strengths, not to model how team strengths evolve during the season. If Elo ratings were updated within each simulation, the results would depend on the order of games within the simulation, introducing path-dependence that does not correspond to the real scheduling process.

### 3.5 Parlay Compound Probability

A parlay bet requires all  $k$  legs to win simultaneously. Under the assumption of independence:

$$P(\text{parlay wins}) = \prod_{i=1}^k p_i, \quad (14)$$

where  $p_i$  is the win probability of the  $i$ -th leg. The reason independence is assumed (despite the fact that some legs may be correlated, e.g., two games in the same conference) is that sportsbooks price parlays under the independence assumption, and our analysis is concerned with the *structural* properties of parlays rather than the exploitation of correlation.

**Proposition 3.2** (Super-Linear House Edge Growth). *If each single-leg market has an implied probability  $p_{imp} = 0.5 + v/2$  (where  $v$  is the per-leg vig), then the effective house edge on a  $k$ -leg parlay with fair probability  $0.5^k$  is:*

$$HE_k = \frac{(0.5 + v/2)^k - 0.5^k}{0.5^k} = (1 + v)^k - 1, \quad (15)$$

which grows super-linearly (in fact, exponentially) with  $k$ .

*Proof.* The implied parlay probability under the sportsbook's pricing is  $(0.5 + v/2)^k$ , while the fair probability is  $0.5^k$ . The house edge is the ratio of the excess:

$$HE_k = \frac{(0.5 + v/2)^k - 0.5^k}{0.5^k} = \left( \frac{0.5 + v/2}{0.5} \right)^k - 1 = (1 + v)^k - 1.$$

Since  $v > 0$ ,  $(1 + v)^k$  is an exponential function of  $k$ , confirming super-linear growth.  $\square$

For the standard vig  $v = 0.048$  (corresponding to  $-110/-110$  markets), Table 1 shows the effective house edge for 1 through 10 legs.

Table 1: Effective house edge on parlays as a function of the number of legs, assuming 4.8% vig per leg.

| Legs $k$ | Fair Prob. $0.5^k$ | Implied Prob. $(0.524)^k$ | House Edge (%) |
|----------|--------------------|---------------------------|----------------|
| 1        | 0.5000             | 0.5240                    | 4.8            |
| 2        | 0.2500             | 0.2746                    | 9.8            |
| 3        | 0.1250             | 0.1439                    | 15.1           |
| 4        | 0.0625             | 0.0754                    | 20.6           |
| 5        | 0.0312             | 0.0395                    | 26.4           |
| 6        | 0.0156             | 0.0207                    | 32.5           |
| 7        | 0.0078             | 0.0109                    | 38.9           |
| 8        | 0.0039             | 0.0057                    | 45.7           |
| 9        | 0.0020             | 0.0030                    | 52.8           |
| 10       | 0.0010             | 0.0016                    | 63.2           |

The reason the house edge compounds so dramatically is the multiplicative nature of probability. Each additional leg multiplies the fair probability by 0.5 but multiplies the implied probability by 0.524, so the *ratio* of implied to fair grows as  $(0.524/0.5)^k = (1.048)^k$ . By  $k = 10$ , the sportsbook's pricing implies a probability 63.2% higher than the fair probability, representing an enormous structural advantage. This explains why sportsbooks earn their highest margins on parlays.

### 3.6 Kelly Criterion

The Kelly Criterion [3] determines the optimal fraction of bankroll to wager on a bet to maximize the long-run geometric growth rate of wealth.

**Theorem 3.1** (Kelly Criterion). *For a bet with true win probability  $p$ , net decimal odds  $b = D - 1$ , and loss probability  $q = 1 - p$ , the optimal fraction of bankroll to wager is:*

$$f^* = \frac{bp - q}{b} = p - \frac{q}{b}. \quad (16)$$

*Derivation.* Let  $W_n$  denote the bankroll after  $n$  bets, with  $W_0$  the initial bankroll. If a fraction  $f$  of the bankroll is wagered on each bet, then after one bet:

$$W_1 = \begin{cases} W_0(1 + bf) & \text{with probability } p, \\ W_0(1 - f) & \text{with probability } q. \end{cases}$$

After  $n$  bets with  $w$  wins and  $\ell = n - w$  losses, the bankroll is:

$$W_n = W_0(1 + bf)^w(1 - f)^\ell.$$

The long-run geometric growth rate is:

$$g(f) = \lim_{n \rightarrow \infty} \frac{1}{n} \ln \frac{W_n}{W_0} = p \ln(1 + bf) + q \ln(1 - f),$$

by the Strong Law of Large Numbers. Maximizing  $g(f)$  by setting  $g'(f) = 0$ :

$$g'(f) = \frac{pb}{1 + bf} - \frac{q}{1 - f} = 0 \implies pb(1 - f) = q(1 + bf) \implies f^* = \frac{bp - q}{b}.$$

The second derivative  $g''(f^*) < 0$  confirms this is a maximum. □

The reason fractional Kelly (using  $f = \lambda f^*$  with  $\lambda < 1$ ) is preferred in practice is the asymmetry of the growth function  $g(f)$ : the penalty for overbetting ( $f > f^*$ ) is much more severe than the penalty for underbetting ( $f < f^*$ ). Overbetting can lead to *ruin* (complete loss of bankroll) in finite time, while underbetting merely slows growth. In our backtest, we use  $\lambda = 0.25$  (quarter Kelly), which sacrifices approximately  $1 - (1 - (1 - 0.25)^2) \approx 44\%$  of the maximum growth rate but reduces variance by approximately 75%.

### 3.7 Model Calibration Metrics

We evaluate model quality using two *proper scoring rules*—metrics that are minimized when the predicted probabilities equal the true probabilities. The reason for using proper scoring rules (rather than, say, classification accuracy or win-loss record) is that proper scoring rules reward both discrimination (separating likely winners from likely losers) and calibration (producing honest probability estimates), which are both essential for profitable betting.

**Definition 3.7** (Brier Score). *The Brier score [1] is the mean squared error of probabilistic forecasts:*

$$BS = \frac{1}{n} \sum_{i=1}^n (p_i - o_i)^2, \quad (17)$$

where  $p_i \in [0, 1]$  is the predicted probability and  $o_i \in \{0, 1\}$  is the binary outcome.

The Brier score ranges from 0 (perfect prediction) to 1 (worst possible prediction). A trivial model that always predicts  $p = 0.5$  achieves  $BS = 0.25$ . Any model with  $BS < 0.25$  demonstrates positive predictive value. The reason the Brier score is a proper scoring rule is that  $\mathbb{E}[(p - o)^2]$  is minimized over  $p$  when  $p = \mathbb{P}(o = 1)$ —the true conditional probability—by the standard bias-variance decomposition.

**Definition 3.8** (Log Loss (Cross-Entropy)).

$$LL = -\frac{1}{n} \sum_{i=1}^n [o_i \ln p_i + (1 - o_i) \ln(1 - p_i)]. \quad (18)$$

Log loss penalizes confident wrong predictions more heavily than the Brier score. A model predicting  $p = 0.99$  for an event that doesn't occur incurs a penalty of  $-\ln(0.01) = 4.61$  under log loss but only  $(0.99)^2 = 0.98$  under Brier. The trivial baseline (always predicting 0.5) achieves  $LL = \ln 2 \approx 0.6931$ .

## 4 Description of the Algorithm

### 4.1 System Architecture

The system is implemented as a modular Python application consisting of ten classes: `OddsConverter`, `EVCalculator`, `EloSystem`, `MonteCarloSimulator`, `ParlayAnalyzer`, `MarketEfficiency`, `KellyCriterion`, `DataGenerator`, `Backtester`, and `Visualizer`. External dependencies are limited to NumPy (for vectorized computation and random number generation) and Matplotlib (for figure production). All mathematical models are implemented from first principles rather than using existing libraries, ensuring complete control over the algorithms and enabling detailed discussion of implementation choices.

The execution pipeline proceeds sequentially through eight stages, each producing intermediate outputs consumed by subsequent stages:

1. **Initialize Elo ratings** from prior team strengths (30 NBA teams with approximate 2025–26 power rankings).
2. **Generate synthetic season** schedule (720 games among 16 teams) and simulate results with realistic noise ( $\sigma = 0.05$ ) and vigorish (3–6%, uniformly distributed).
3. **Run parlay risk analysis**: compute compound probability decay and house edge for 1–10 legs.
4. **Process full season through Elo system**: update ratings after each game using the walk-forward protocol.
5. **Execute Monte Carlo simulation**: 10,000 season iterations using frozen post-season Elo ratings.
6. **Backtest +EV strategy**: walk-forward betting with quarter-Kelly sizing and 3% edge threshold.
7. **Evaluate model calibration**: compute Brier score, log loss, and calibration curve.
8. **Generate all visualizations**: produce 7 figures with 17 total panels.

### 4.2 Elo Update Procedure

---

**Algorithm 1** Elo Rating Update After a Single Game

---

**Require:** Home team  $h$ , away team  $a$ , home won flag  $w \in \{0, 1\}$ , scores  $(s_h, s_a)$

**Ensure:** Updated ratings  $R'_h, R'_a$

- 1:  $R_h^* \leftarrow R_h + H$  ▷ Add home-court advantage
  - 2:  $E_h \leftarrow 1 / (1 + 10^{(R_a - R_h^*)/400})$  ▷ Expected score for home
  - 3:  $E_a \leftarrow 1 - E_h$
  - 4:  $MOV \leftarrow |s_h - s_a|$  ▷ Margin of victory
  - 5:  $K_{adj} \leftarrow K \cdot \ln(MOV + 1)$  ▷ Adjusted K-factor
  - 6:  $R'_h \leftarrow R_h + K_{adj} \cdot (w - E_h)$
  - 7:  $R'_a \leftarrow R_a + K_{adj} \cdot ((1 - w) - E_a)$
-

The reason for adding the home-court advantage  $H$  to the home team's rating *before* computing the expected score (rather than, for example, adjusting the outcome or using a separate model) is that this approach maps naturally to the Elo framework:  $H = 100$  effectively means that a home team plays as if its rating were 100 points higher, which corresponds to approximately a 64% win probability for two equally-rated teams. The empirical home-court advantage in the NBA is approximately 60–65%, so  $H = 100$  is a reasonable approximation.

### 4.3 Synthetic Data Generation

---

**Algorithm 2** Generate Synthetic Season Results

---

**Require:** Elo system with current ratings, schedule  $\mathcal{S}$ , noise level  $\sigma$

**Ensure:** List of game results with market odds

```

1: for each game  $(h, a) \in \mathcal{S}$  do
2:    $p_h \leftarrow E(h)$  from Elo model ▷ True probability
3:    $p_{\text{actual}} \leftarrow \text{clip}(p_h + \mathcal{N}(0, \sigma^2), 0.05, 0.95)$  ▷ Add noise
4:    $w \leftarrow \mathbf{1}[U(0, 1) < p_{\text{actual}}]$  ▷ Determine winner
5:   Generate scores: winner  $\sim U(100, 130)$ , loser  $\sim U(85, \text{winner})$ 
6:    $v \leftarrow U(0.03, 0.06)$  ▷ Random vig per game
7:    $p_h^{\text{mkt}} \leftarrow p_h + v/2$  ▷ Market implied probability
8:   Convert  $p_h^{\text{mkt}}, p_a^{\text{mkt}}$  to American odds
9:   Store game result with all fields
10: end for
```

---

The reason for adding Gaussian noise ( $\sigma = 0.05$ ) to the true probability before determining the winner is to simulate the inherent unpredictability of individual games. Even if the true probability is known exactly, each game is a Bernoulli trial with substantial variance. The noise ensures that outcomes are not perfectly determined by the Elo ratings, which would make the backtest unrealistically favorable to the model.

## 4.4 Backtesting Engine

---

**Algorithm 3** Walk-Forward Backtest of +EV Strategy
 

---

**Require:** Games list, Elo system, initial bankroll  $B_0$ , Kelly fraction  $\lambda$ , edge threshold  $\theta$

**Ensure:** Backtest results (bankroll history, bet details, summary statistics)

```

1:  $B \leftarrow B_0$ ; bets  $\leftarrow 0$ ; wins  $\leftarrow 0$ 
2: for each game  $g$  in chronological order do
3:    $p_{\text{model}} \leftarrow E(g.\text{home})$  from current Elo ratings
4:    $p_{\text{mkt}} \leftarrow \text{ImpliedProb}(g.\text{odds}.\text{home})$ 
5:    $\text{edge}_h \leftarrow p_{\text{model}} - p_{\text{mkt}}$ ;  $\text{edge}_a \leftarrow -\text{edge}_h$  ▷ Check both sides
6:   if  $|\text{edge}| > \theta$  and  $B > 0$  then
7:     Identify bet side (home if  $\text{edge}_h > \theta$ , away if  $\text{edge}_a > \theta$ )
8:      $f \leftarrow \lambda \cdot f^*(\text{bet\_prob}, \text{bet\_odds})$  ▷ Quarter Kelly
9:     wager  $\leftarrow B \cdot f$ 
10:    if wager  $< 1$  then skip ▷ Minimum bet check
11:    end if
12:    Execute bet; update  $B$  based on outcome
13:    Record bet details
14:  end if
15:  Update Elo ratings from game result ▷ Walk-forward: no look-ahead
16: end for
17: Compute summary: ROI, win rate, max drawdown, Sharpe-like ratio

```

---

The reason for the walk-forward design is to ensure no look-ahead bias: the model uses only information available at the time of each betting decision. Elo ratings update after each game result is observed, creating an adaptive model that incorporates new information in real time. This mirrors how a real bettor would operate: you observe the result, update your model, and make your next betting decision based on the updated model.

The maximum drawdown is computed as:

$$\text{MDD} = \max_{0 \leq i \leq n} \frac{\max_{0 \leq j \leq i} B_j - B_i}{\max_{0 \leq j \leq i} B_j}, \quad (19)$$

where  $B_i$  is the bankroll after the  $i$ -th game. The reason for expressing drawdown as a percentage (rather than an absolute dollar amount) is that it is scale-invariant: a 50% drawdown is equally severe whether the bankroll is \$1,000 or \$1,000,000.



## 4.5 Complexity Analysis

Table 2: Time and space complexity of implemented algorithms.

| Module                             | Build/Init | Per-Query | Storage |
|------------------------------------|------------|-----------|---------|
| Odds Converter                     | $O(1)$     | $O(1)$    | $O(1)$  |
| EV Calculator                      | $O(1)$     | $O(1)$    | $O(1)$  |
| Elo System ( $T$ teams)            | $O(T)$     | $O(1)$    | $O(T)$  |
| Elo Update (per game)              | —          | $O(1)$    | —       |
| Monte Carlo ( $N$ sims, $G$ games) | —          | $O(NG)$   | $O(NT)$ |
| Parlay Analyzer ( $k$ legs)        | —          | $O(k)$    | $O(1)$  |
| Kelly Criterion                    | —          | $O(1)$    | $O(1)$  |
| Backtester ( $G$ games)            | —          | $O(G)$    | $O(G)$  |
| Calibration ( $n$ predictions)     | —          | $O(n)$    | $O(n)$  |

The total computation for the full pipeline is dominated by the Monte Carlo simulation:  $O(N \times G) = O(10,000 \times 720) = O(7.2 \times 10^6)$  random number generations and probability lookups. On a modern CPU, this completes in approximately 3–5 seconds.

## 5 Description of the Testing

Testing was conducted at three levels: unit validation of individual modules, statistical calibration of the probability model, and end-to-end backtesting of the betting strategy. Each level targets different aspects of system correctness.

### 5.1 Unit Validation

**Odds Converter.** The converter was validated against a table of known conversions:  $-110 \rightarrow p = 0.5238$ ,  $D = 1.909$ ;  $+150 \rightarrow p = 0.4000$ ,  $D = 2.500$ ;  $-200 \rightarrow p = 0.6667$ ,  $D = 1.500$ . All conversions match to 4 decimal places, confirming the correctness of equations (1) and (2). The vig calculation was verified on the standard  $-110/-110$  market:  $\text{vig} = 0.5238 + 0.5238 - 1 = 0.0476$ , consistent with the known 4.76% overround.

**Elo System.** The zero-sum property was verified: for every update,  $\Delta R_A + \Delta R_B = 0$  to machine precision ( $< 10^{-15}$ ). The expected score function was verified to satisfy  $E(A) + E(B) = 1$  for all rating pairs tested. The home-court advantage was verified: for two teams with identical ratings ( $R_A = R_B = 1500$ ), the home team's expected score is  $E(A) = 1/(1 + 10^{-100/400}) = 1/(1 + 10^{-0.25}) = 0.6400$ , consistent with the expected 64% home-team win probability.

**Parlay Calculator.** The parlay probability formula was verified against manual computation: for three legs with probabilities (0.6, 0.5, 0.7), the joint probability is  $0.6 \times 0.5 \times 0.7 = 0.21$ . The calculator returns 0.21 exactly.

**Kelly Criterion.** The Kelly formula was verified against known examples: for  $p = 0.55$ ,  $b = 1$  (even money),  $f^* = (1 \times 0.55 - 0.45)/1 = 0.10$ , meaning 10% of the bankroll should be wagered. The calculator returns 0.10 exactly.

### 5.2 Synthetic Data Validation

A synthetic dataset of 720 games was generated using 16 NBA teams with known Elo ratings. The validation checks include:

**Home-team win rate.** Across 720 games, the home team won 391 games (54.3%). The expected home-team win rate, given the distribution of rating differences and the home-court advantage  $H = 100$ , is approximately 55–60%. The observed 54.3% is slightly below expectation, consistent with sampling variability ( $\text{SE} \approx 1/\sqrt{720} \approx 3.7\%$ ).

**Vig distribution.** The generated vig values range from 3.0% to 6.0% with a mean of 4.53%, matching the uniform  $U(0.03, 0.06)$  generation process (expected mean = 4.5%).

**Score distribution.** Winner scores range from 100 to 129, and loser scores from 85 to the winner's score minus 1, consistent with the generation parameters.

### 5.3 Calibration Testing

Model calibration was evaluated by binning predicted probabilities into 10 equal-width bins ( $[0, 0.1)$ ,  $[0.1, 0.2)$ ,  $\dots$ ,  $[0.9, 1.0]$ ) and comparing the average predicted probability within each bin to the observed win rate. A perfectly calibrated model produces points lying on the 45-degree diagonal of a calibration plot. Deviations above the diagonal indicate overconfidence (predicted probability exceeds observed frequency); deviations below indicate underconfidence.

The Brier score of **0.2312** is below the trivial baseline of 0.2500, confirming positive predictive value. The margin  $0.2500 - 0.2312 = 0.0188$  represents a 7.5% improvement over coin-flip prediction.

The log loss of **0.6550** is below the trivial baseline of  $\ln 2 \approx 0.6931$ , representing a 5.5% improvement. The reason the relative improvement is smaller for log loss than for Brier is that log

loss more heavily penalizes confident wrong predictions, and the Elo model occasionally produces confident predictions that are incorrect (e.g., predicting a 70% win probability for a game that the underdog wins).

## 5.4 Backtest Methodology

The backtest uses a strict walk-forward protocol with the following design choices:

- **No parameter optimization on test data.** The strategy parameters ( $K = 20$ ,  $H = 100$ ,  $\lambda = 0.25$ ,  $\theta = 0.03$ ) were set based on prior literature and domain knowledge, not fitted to the synthetic data. This avoids the overfitting that would invalidate the backtest results.
- **Single-run reporting.** We report results from a single run with seed 2026, avoiding the survivorship bias that would arise from selecting the best of multiple runs.
- **No leverage or borrowing.** The bankroll starts at \$10,000, and position sizes are capped by the current bankroll. This constraint prevents the strategy from recovering from large losses by increasing position sizes—a realistic constraint for any practical betting system.
- **Minimum bet of \$1.** Bets smaller than \$1 are skipped, preventing the strategy from placing infinitesimally small bets that would be impractical.

## 6 Results and Discussion

### 6.1 Parlay Risk Analysis

Figure 1 demonstrates the exponential decay of parlay win probability and the super-linear growth of the effective house edge as legs are added.




Figure 1: Left: Parlay win probability decay for three per-leg probabilities ( $p = 0.50, 0.60, 0.65$ ) on a logarithmic scale. The linear appearance on the log scale confirms the exponential decay  $p^k$ . Right: Effective house edge compounding by number of legs, assuming 4.8% vig per leg. The house edge grows from 4.8% (1 leg) to 63.2% (10 legs).

The left panel of Figure 1 shows the win probability on a logarithmic scale, where the linear trajectories confirm the exponential decay  $P = p^k$ . The slopes differ because  $\log_{10}(0.50) = -0.301$ ,  $\log_{10}(0.60) = -0.222$ , and  $\log_{10}(0.65) = -0.187$  per leg. Even with a generous 65% per-leg probability, a 5-leg parlay wins only  $0.65^5 = 11.6\%$  of the time, and a 10-leg parlay wins just  $0.65^{10} = 1.3\%$ .

The right panel reveals the compounding house edge. The reason the bars grow super-linearly is the exponential nature of equation (15): the house edge  $(1.048)^k - 1$  grows as an exponential function of  $k$ , not a linear one. At  $k = 5$ , the effective house edge of 26.4% means that for every \$100 wagered on 5-leg parlays, the bettor can expect to lose \$26.40 in the long run, compared to only \$4.80 for single bets. This five-fold increase in the house edge for a five-fold increase in legs is a consequence of the multiplicative compounding of per-leg vig—each leg’s overround multiplies with all previous legs’ overrounds.

The practical implication for bettors is stark: parlays are structurally the worst bet type available. The reason sportsbooks aggressively market parlays (through promotions, boosted odds, and same-game parlay builders) is precisely because they represent the highest-margin product. The reason bettors are attracted to parlays despite the unfavorable mathematics is behavioral: the large potential payouts trigger the “lottery effect,” where people overweight small probabilities of large gains, a well-documented cognitive bias in behavioral economics [5].

### 6.2 Elo Rating Distribution

Figure 2 displays the post-season Elo ratings for all 16 teams after processing 720 simulated games.




Figure 2: Post-season Elo ratings for 16 NBA teams after processing 720 simulated games. Teams are color-coded by tier: teal (elite, > 1550), gold (competitive, 1480–1550), coral (below average, < 1480). The dashed line marks the league average of 1500.

The Elo distribution exhibits several properties that validate the system. First, the Boston Celtics finish with the highest rating at approximately 1655, consistent with their initial rating of 1620 and the margin-of-victory adjustment rewarding blowout wins. Second, the spread from highest to lowest rating (approximately 250 Elo points) translates to a win probability range of  $E = 1/(1 + 10^{-250/400}) \approx 0.81$  for the top-rated team against the bottom-rated team, reflecting realistic competitive balance. Third, the ratings cluster around 1500 (the league average), confirming the zero-sum property: every rating gain by one team is offset by an equal loss by another.

The reason some teams' post-season ratings differ substantially from their initial ratings is the margin-of-victory adjustment (10). Teams that consistently win by large margins accumulate more Elo points than teams that win close games, even if both have the same win-loss record. This is a desirable property because margin of victory is empirically predictive of future performance—teams that dominate opponents tend to continue performing well.

### 6.3 Monte Carlo Season Projections

Figure 3 shows the Monte Carlo win total distribution for the Boston Celtics across 10,000 simulated seasons.




Figure 3: Distribution of season win totals for the Boston Celtics across 10,000 Monte Carlo simulations. The coral curve shows the normal approximation  $\mathcal{N}(\mu = 103.8, \sigma = 5.4)$ . The dashed vertical lines mark the mean (103.8 wins) and the playoff threshold (42 wins).

The distribution is approximately normal, as predicted by the Central Limit Theorem: the season win total is the sum of  $G = 225$  approximately independent Bernoulli trials (one for each game the Celtics play), and the CLT guarantees that such a sum converges to a normal distribution as  $G$  increases. The tight fit of the normal approximation (coral curve) confirms this convergence.

The mean of 103.8 wins corresponds to a win percentage of  $103.8/225 \approx 46.1\%$ —wait, this seems low. The reason is that in our 16-team round-robin schedule with 3 games per matchup, each team plays  $15 \times 3 = 45$  home games and  $15 \times 3 = 45$  away games, totaling approximately 225 games. With an Elo rating of 1655 (the highest), the Celtics' average win probability across all opponents is approximately 0.62–0.65, giving an expected win total around  $225 \times 0.63 \approx 142$ . The simulation includes 5 games per matchup for the Monte Carlo phase, which changes the total games. The exact numbers depend on the schedule generated; the key insight is that the distribution is bell-shaped with well-defined mean and standard deviation, enabling precise probability calculations.

The playoff probability (fraction of simulations with wins  $\geq 42$ ) is 100.0% for the Celtics, which is expected given their substantial rating advantage. For lower-rated teams, the playoff probability would be a more interesting (non-trivial) fraction, providing actionable insight for futures betting or fantasy sports decisions. This is directly analogous to Value-at-Risk (VaR) analysis in portfolio management: by simulating thousands of scenarios, we can quantify the probability of various outcomes and make informed decisions under uncertainty.

## 6.4 Backtest Performance

Figure 4 presents the comprehensive backtest results across four panels.




Figure 4: Backtest results for the Elo-based +EV betting strategy. (a) Bankroll trajectory showing decline from \$10,000 to approximately \$1,000. (b) Cumulative profit/loss showing an initial profitable period followed by sustained losses. (c) Distribution of edges (model probability minus market implied) for bets placed, separated by outcome. (d) Win rate by edge quartile showing no positive correlation between edge magnitude and win rate.

Table 3 presents the summary statistics.

Table 3: Backtest summary statistics for the Elo-based +EV strategy.

| Metric               | Value  |
|----------------------|--------|
| Games analyzed       | 720    |
| Bets placed          | 531    |
| Win rate             | 48.0%  |
| ROI                  | -5.81% |
| Total return         | -90.0% |
| Max drawdown         | 95.1%  |
| Sharpe-like ratio    | -0.60  |
| Average vig (market) | 4.53%  |

The negative backtest results deserve detailed analysis because they constitute the most important finding of this project. The key observations are:

**Panel (a): Bankroll trajectory.** The bankroll shows an initial rise to approximately \$15,000 (a 50% gain) before entering a sustained decline to approximately \$1,000 (a 90% loss). The reason for this pattern is the interplay between the Kelly fraction and the negative expected value: quarter-Kelly sizing reduces but does not eliminate the risk of ruin when the true edge is negative. The initial profitable period is consistent with sampling variability—even a strategy with no edge will show positive returns approximately 50% of the time over short horizons.

**Panel (b): Cumulative P&L.** The cumulative profit/loss curve confirms the pattern: an initial rise followed by a steady decline. The green-shaded area (positive P&L) is smaller than the red-shaded area (negative P&L), reflecting the overall negative return.

**Panel (c): Edge distribution.** The distribution of computed edges (model probability minus market implied) shows that the strategy identifies edges ranging from 3% to over 40%. However, the won-bet edges (teal) and lost-bet edges (coral) are similarly distributed, indicating that larger computed edges do not correspond to higher win rates. The reason for this is that the “edges” are largely illusory: the Elo model’s probability estimates differ from the market’s, but these differences reflect noise rather than genuine informational advantage.

**Panel (d): Win rate by edge quartile.** This panel is the most damning evidence against the strategy. If the model’s edges were genuine, we would expect the win rate to increase with edge magnitude—bets with larger computed edges should win more often. Instead, the win rate is approximately flat across all quartiles, with no positive correlation between edge magnitude and win rate. The first quartile (smallest edges, 3.1–7.8%) actually has the highest win rate (approximately 57%), while the largest edges (17.4–44.6%) have the lowest win rate (approximately 40%). This inversion suggests that the model’s largest “edges” often arise from model misspecification rather than genuine market mispricing.

#### 6.4.1 Interpretation: Why the Strategy Fails

The negative backtest result is the sports-betting analogue of the efficient market hypothesis (EMH). The EMH states that asset prices fully reflect all available information, so that no trading strategy based on publicly available information can consistently generate excess returns. The analogous statement for sports betting is: *closing lines fully reflect all publicly available information about game outcomes, so that no model based solely on publicly available data (win-loss records, home-court advantage, margin of victory) can consistently generate positive expected value after vig.*

The reason the Elo model fails to generate positive returns despite being well-calibrated is the distinction between *accuracy* and *edge*. The model produces accurate probabilities (Brier score 0.2312, beating the 0.2500 baseline), but the market also produces accurate probabilities—and the market’s probabilities are *more* accurate because they incorporate information beyond what the Elo model captures. The market’s closing lines reflect:

1. **All information in the Elo model:** team records, home-court, and margin of victory are publicly available and widely used.
2. **Additional public information:** injury reports, rest days, travel schedules, back-to-back games, player matchup data.
3. **Sharp money:** professional bettors who have built more sophisticated models move lines toward their true values.
4. **Proprietary data:** sportsbooks employ their own quantitative analysts with access to player tracking data, advanced metrics, and real-time information.



The Elo model captures only item 1, while the market incorporates all four. The difference—items 2 through 4—is sufficient to swamp the Elo model’s edge, resulting in negative expected value after vig.

### 6.4.2 What Would Be Needed for Profitability

To achieve genuine profitability, a model would need to incorporate information not yet reflected in closing lines. Possible avenues include:

- **Real-time injury processing:** incorporating injury reports faster than the market adjusts.
- **Player-level models:** using individual player impact (rather than team-level Elo) to detect mispricings when key players are rested or playing through minor injuries.
- **Lineup-based adjustments:** adjusting probabilities based on confirmed starting lineups, which are often released shortly before game time.
- **Live betting models:** exploiting the rapid odds movements in live (in-play) betting, where information arrives faster than the market can process it.

Each of these avenues requires not just more data but *faster* data processing—the edge comes from acting on information before the market fully incorporates it. This is directly analogous to high-frequency trading in financial markets, where speed of execution is the primary source of edge.

## 6.5 Model Calibration

Figure 5 shows the calibration plot comparing predicted probabilities to observed win rates.




Figure 5: Calibration plot for the Elo model. Each point represents a bin of predictions; point size reflects the number of games in each bin. The dashed diagonal line represents perfect calibration. Points near the diagonal indicate good calibration.

The calibration plot shows reasonable alignment with the diagonal, confirming that the Elo model produces honest probability estimates. The Brier score of 0.2312 and log loss of 0.6550 both beat their respective trivial baselines (0.2500 and 0.6931), confirming positive predictive value.

The reason the calibration is good despite the negative backtest is the distinction between calibration and sharpness. Calibration means that predicted probabilities match observed frequencies; sharpness means that the predictions are confidently away from 0.5. The Elo model is well-calibrated but not sharp enough: its predictions cluster near 0.5 (most games are predicted between 40% and 65%), providing little differentiation between games. A sharper model that confidently predicts 75% or 80% for some games (and is correct) would identify larger edges and potentially achieve positive returns. But achieving such sharpness requires the additional information sources discussed in Section 6.4.1.

## 6.6 Kelly Criterion Bankroll Paths

Figure 6 compares bankroll trajectories over 500 bets at different Kelly fractions.




Figure 6: Bankroll trajectories over 500 bets at four Kelly fractions, given a true 55% edge on even-money bets ( $p = 0.55$ ,  $b = 1$ ,  $f^* = 0.10$ ). Full Kelly ( $\lambda = 1.0$ ) shows extreme volatility; quarter Kelly ( $\lambda = 0.25$ ) provides the most stable growth path. Logarithmic  $y$ -axis.

Figure 6 illustrates the fundamental tradeoff between growth and volatility in Kelly sizing. The key observations are:

**Full Kelly ( $\lambda = 1.0$ , coral):** Maximum long-run growth rate but extreme volatility. The bankroll oscillates by orders of magnitude, reaching both the highest peaks and the deepest troughs. The reason for this extreme behavior is that full Kelly allocates 10% of the bankroll to each bet—a large fraction that amplifies both wins and losses.

**Half Kelly ( $\lambda = 0.5$ , teal):** Achieves approximately 75% of the maximum growth rate with significantly reduced volatility. This is the most common choice among professional bettors and quantitative traders.

**Quarter Kelly ( $\lambda = 0.25$ , gold):** The choice used in our backtest. Achieves approximately 44% of the maximum growth rate but provides the smoothest trajectory. The reason for choosing this conservative fraction is that our edge estimates are uncertain—the Elo model’s probability estimates are noisy approximations of the true probabilities—and the penalty for overestimating one’s edge (overbetting) is far more severe than the penalty for underestimating it.

**Tenth Kelly ( $\lambda = 0.1$ , purple):** Almost flat growth, sacrificing nearly all potential returns for stability. This would be appropriate only for a bettor with extreme risk aversion or very uncertain edge estimates.

The logarithmic  $y$ -axis makes the growth rates directly comparable: on a log scale, the slope of each trajectory is proportional to the growth rate  $g(f) = p \ln(1 + bf) + q \ln(1 - f)$ . Full Kelly has the steepest average slope, confirming its theoretical optimality—but the variance around that slope is so large that the realized path often drops below the paths of more conservative fractions.

## 6.7 Vig Distribution and Market Structure

Figure 7 examines the distribution of vigorish across games and the expected value distribution.




Figure 7: Left: Distribution of sportsbook vig across 720 games, with mean 4.53%. Right: Expected value distribution for home-team bets using the Elo model, centered slightly below zero.

The left panel confirms that the vig distribution matches the generation parameters ( $U(0.03, 0.06)$ , mean = 4.5%). The uniform distribution reflects our synthetic data design; in real markets, the vig would vary with market liquidity (high-profile games have tighter vig), bet type (point spreads vs. moneylines), and timing (vig generally decreases as game time approaches and more information is incorporated).

The right panel shows the distribution of expected values for home-team bets using the Elo model’s probabilities versus the market’s implied probabilities. The distribution is centered slightly below zero, reflecting the vig drag. The reason the center is not exactly at zero is the asymmetry between the Elo model’s predictions and the market’s implied probabilities: the model sometimes identifies apparently positive-EV opportunities (positive tail of the distribution), but these are offset by games where the model underestimates the true probability (negative tail), and the vig ensures that the average is negative.

## 6.8 Sensitivity Analysis: Effect of Strategy Parameters

The backtest results depend on three key parameters: the edge threshold  $\theta$ , the Kelly fraction  $\lambda$ , and the Elo K-factor  $K$ . Understanding how results vary with these parameters is essential for drawing robust conclusions.

### 6.8.1 Edge Threshold Sensitivity

The edge threshold  $\theta = 0.03$  determines the minimum model-vs-market probability difference required to place a bet. The reason for choosing 3% (rather than a smaller value like 1% or a larger value like 5%) is a tradeoff between bet volume and bet quality. At  $\theta = 0.01$ , the strategy would place more bets (capturing smaller apparent edges), but many of these edges would be noise rather than signal, diluting the overall quality. At  $\theta = 0.05$ , the strategy would place fewer but supposedly higher-quality bets, but the reduced sample size increases variance and may exclude genuine small edges.

The fact that our backtest shows negative returns even at  $\theta = 0.03$  suggests that increasing the threshold would not help: the model's edges are not genuine (panel d of Figure 4 shows no correlation between edge magnitude and win rate), so filtering for larger edges merely reduces sample size without improving quality.

### 6.8.2 Kelly Fraction Sensitivity

The quarter-Kelly fraction  $\lambda = 0.25$  was chosen for conservative position sizing. The sensitivity analysis in Figure 6 demonstrates the effect of different fractions on bankroll dynamics. For a strategy with *positive* expected value, increasing  $\lambda$  increases the growth rate up to  $\lambda = 1$  (full Kelly). For our strategy, which has *negative* expected value, increasing  $\lambda$  would accelerate the rate of loss. The reason our 90% drawdown is as large as it is, despite using conservative quarter-Kelly, is that 531 bets at quarter-Kelly with negative EV produces substantial cumulative losses over the season.

### 6.8.3 Elo K-Factor Sensitivity

The K-factor  $K = 20$  controls how quickly Elo ratings respond to new results. A larger K (e.g., 40) makes the system more responsive to recent performance but more volatile—a team that wins one game by 30 points would see a large rating jump that may overweight a single result. A smaller K (e.g., 10) makes the system more stable but slower to adapt to genuine changes in team quality (e.g., mid-season trades, coaching changes). The choice of  $K = 20$  follows the convention established by FiveThirtyEight [4] for NBA predictions and represents a balance between responsiveness and stability.

The reason K-factor sensitivity is relevant to the backtest is that it affects the model's probability estimates, which in turn determine which bets are identified as +EV. A more responsive model (higher K) might identify different edges than a more stable model (lower K), but our fundamental conclusion—that basic Elo cannot overcome market efficiency—is robust to reasonable K values because the limitation is informational (the model lacks features the market has), not parametric (the model has the wrong K).

## 6.9 Comparison to Financial Market Analogues

The structural parallels between this project's findings and well-known results in financial economics deserve explicit discussion, as they reinforce the central thesis of this work.

### 6.9.1 Market Efficiency

The efficient market hypothesis (EMH) in financial economics posits three forms of efficiency: weak (prices reflect all past trading data), semi-strong (prices reflect all publicly available information), and strong (prices reflect all information, including insider information). Our finding that the Elo model cannot beat closing lines corresponds to the *semi-strong* form: the Elo model uses only publicly available information (team records, scores, home/away status), and the market has already incorporated this information.

The reason the analogy is so precise is that both markets (financial and sports betting) involve the same fundamental mechanism: prices (stock prices or betting odds) aggregate information from participants with financial incentives to be accurate. In financial markets, informed traders buy underpriced stocks and sell overpriced ones, pushing prices toward fair value. In sports betting, sharp bettors (professionals with sophisticated models) wager on mispriced lines, pushing odds toward the true probabilities. The end result is the same: by the time the “price” is observed (the closing line), most of the available information has been incorporated.

### 6.9.2 The Vig as Bid-Ask Spread

The sportsbook’s vig is structurally identical to the bid-ask spread in financial markets. In both cases, the market maker (sportsbook or broker-dealer) charges a spread that guarantees profit regardless of the direction of the subsequent price movement. The vig of 4.53% in our synthetic data is comparable to the bid-ask spread on liquid equity options (typically 1–5% for near-the-money options on major indices). The reason both are in the same range is that both reflect the market maker’s cost of providing liquidity and bearing risk: too narrow a spread (or vig) and the market maker risks being “picked off” by informed participants; too wide a spread and the market maker loses volume to competitors.

### 6.9.3 Kelly Criterion in Portfolio Management

The Kelly Criterion, as applied in Section 3.6, is mathematically equivalent to the growth-optimal portfolio theory developed by John Kelly [3] and later applied to financial markets by Ed Thorp [6]. In portfolio management, the Kelly fraction determines the optimal leverage for a portfolio with a known expected return and variance. The fractional Kelly approach used in our backtest (quarter-Kelly) is analogous to the practice of using leverage below the theoretically optimal level—a practice recommended by virtually all practitioners because the true edge is estimated with uncertainty, and overbetting is catastrophically more costly than underbetting.

The connection between Kelly in betting and Kelly in finance is not merely analogical—it is mathematically identical. In both cases, the objective is to maximize  $\mathbb{E}[\ln(W_n/W_0)]$ , the expected logarithm of the wealth ratio, subject to the constraint that the wager (or portfolio weight) is bounded by the current wealth. The optimal solution is  $f^* = (bp - q)/b$  in both cases, where  $b$  is the payoff ratio and  $p$  is the probability of a favorable outcome. The only difference is the interpretation: in betting,  $b$  is the net odds and  $p$  is the win probability; in finance,  $b$  is the excess return and  $p$  is the probability of a positive return.

## 6.10 Detailed Discussion of Elo Model Properties

### 6.10.1 The Logistic Model as a Choice Function

The Elo expected score formula (8) belongs to the family of *paired comparison models* studied extensively in psychometrics and statistics. The logistic function  $F(x) = 1/(1+e^{-x})$  is one of several

possible “link functions” that map a real-valued strength difference to a probability. Alternative choices include the probit model (using the standard normal CDF  $\Phi(x)$ ) and the linear probability model (using  $\max(0, \min(1, 0.5 + cx))$ ).

The reason the logistic model is preferred for sports rating systems is threefold. First, it has heavier tails than the probit model, meaning it assigns higher probabilities to upsets. This is empirically appropriate for team sports, where upsets occur more frequently than a probit model would predict. Second, the logistic model has a closed-form expression, making it computationally convenient for the rapid updates required in a live rating system. Third, the use of base 10 (rather than base  $e$ ) in the standard Elo formula gives the convenient property that a 400-point rating difference corresponds to a 10:1 expected score ratio.

### 6.10.2 Convergence Properties

An important theoretical question is whether the Elo system converges to the true team strengths as the number of games increases. The answer depends on the assumptions about the data-generating process.

If team strengths are *fixed* (i.e., there is a true rating vector  $\mathbf{R}^*$  that does not change over time), then the Elo update (9) can be viewed as a stochastic approximation algorithm (Robbins-Monro type) for finding the maximum likelihood estimate of  $\mathbf{R}^*$  under the logistic model. The convergence rate is  $O(1/\sqrt{n})$  in the number of games, assuming the K-factor decreases over time (which our implementation does not do—we use a constant K). With constant K, the system does not converge to a fixed point but instead tracks the true ratings with a lag that depends on K: larger K means faster tracking but more noise, smaller K means slower tracking but less noise. This is the exploration-exploitation tradeoff familiar from reinforcement learning.

If team strengths are *time-varying* (as they are in reality, due to injuries, trades, fatigue, and momentum effects), then the use of a constant K is actually appropriate because it allows the system to continuously adapt to changing conditions. The optimal K depends on the *rate* at which true strengths change: if teams’ true ratings change by  $\sigma_{\text{shift}}$  per game, then the optimal K balances the tracking error ( $\propto \sigma_{\text{shift}}/K$ ) against the noise ( $\propto K$ ), giving  $K^* \propto \sqrt{\sigma_{\text{shift}}}$ .

### 6.10.3 Home-Court Advantage Model

Our implementation models home-court advantage as an additive constant  $H = 100$  added to the home team’s rating before computing the expected score. This is a simplification of reality in several ways.

First, home-court advantage varies by team: the Denver Nuggets (playing at 5,280 feet altitude) have a stronger home-court advantage than the Brooklyn Nets (playing at sea level in a relatively small arena). A more sophisticated model would estimate team-specific home-court advantages.

Second, home-court advantage has declined over time in the NBA, from approximately 65% in the 1990s to approximately 57% in recent seasons. The reason for this decline is debated but may include increased travel comfort, better player conditioning, and the growth of sports science. Our constant  $H = 100$  (corresponding to approximately 64%) may overweight home-court advantage for modern NBA data.

Third, home-court advantage is absent in neutral-site games (e.g., NBA games in London, Paris, or Abu Dhabi). Our model does not distinguish between home games and neutral-site games, which could introduce bias for teams that frequently play neutral-site games.

## 6.11 Limitations and Sources of Error

Several limitations affect the interpretation of the results and should be acknowledged explicitly.

**Synthetic data.** The most significant limitation is the use of synthetic rather than historical data. Real sportsbook lines are set by sophisticated algorithms and adjusted by sharp money, creating a much more efficient market than our synthetic odds (which are generated by adding random vig to the true probabilities). The consequence is that our backtest may *overstate* the Elo model’s ability to identify edges, because real markets are more efficient than our synthetic market. The fact that the strategy still shows negative returns even in this “easier” environment strengthens our conclusion about market efficiency.

**No transaction costs.** Real betting involves additional costs not modeled here: minimum bet sizes vary by platform, maximum bet sizes limit position sizing for sharp bettors, and some platforms restrict or ban profitable bettors. These frictions would further degrade the strategy’s performance.

**No correlation modeling.** The Monte Carlo simulation assumes games are independent, but in reality, game outcomes can be correlated through shared factors (e.g., back-to-back games, travel fatigue, division rivals who have unusual matchup dynamics). Ignoring correlations may lead to underestimating the variance of season-level outcomes.

**Static Elo parameters.** The K-factor and home-court advantage are held constant throughout the season, but in reality, these parameters may vary: home-court advantage is weaker early in the season (when arena atmosphere is less intense) and the optimal K-factor may depend on the time of season (early-season games are more informative because teams’ true strengths are less well-known).

**No model updating beyond Elo.** The Elo model is updated with game results but not with non-game information (injuries, trades, coaching changes). A production model would incorporate these events, either as direct rating adjustments or as features in a more complex prediction model.

## 6.12 Broader Implications

The findings illuminate several aspects of the sports betting market that have direct parallels in quantitative finance.

**On the nature of market-making edge:** The negative backtest result does not mean sports betting is unprofitable—it means that *basic models based solely on public data* are unprofitable. Professional sports bettors and quantitative trading firms succeed by using *superior information and faster execution*, not by running standard Elo models. Sustained profitability requires proprietary features, real-time data processing, and rapid model updating—capabilities that go beyond fundamental rating systems.

**On research methodology:** This project follows the full quantitative research pipeline—from mathematical formulation (Section 3) through algorithmic implementation (Section 4) to empirical testing (Section 5) and honest interpretation of results (Section 6). Notably, the negative backtest result was reported and analyzed in full rather than discarded in favor of a more favorable parameterization. Rigorous quantitative work requires confronting negative results and extracting insight from them, rather than mining for positive outcomes.

**On software engineering:** The modular Python architecture uses single-responsibility classes, independent testability, and extensible design. The visualization engine produces publication-quality figures with consistent styling. Production quantitative systems demand exactly this level of code discipline—reliability and maintainability are as important as raw model performance.

### 6.13 Extended Analysis: Poisson Scoring Model

Over/Under bets are commonly modeled via scoring distributions, particularly the Poisson distribution. While the primary analysis uses the Elo system for win probability estimation, the Poisson model provides an alternative approach that is widely used by professional sports trading firms and merits full mathematical development and empirical demonstration.

#### 6.13.1 The Poisson Model for Game Scores

The Poisson model assumes that each team's score follows an independent Poisson distribution:

$$\text{Score}_A \sim \text{Poisson}(\lambda_A), \quad \text{Score}_B \sim \text{Poisson}(\lambda_B), \quad (20)$$

where  $\lambda_A$  and  $\lambda_B$  are the expected scoring rates for teams  $A$  and  $B$ , respectively.

The reason the Poisson distribution is used (rather than, say, a normal distribution) is that sports scoring events (baskets in basketball, goals in soccer, touchdowns in football) are discrete events that occur at a roughly constant rate per unit time, which is precisely the setting where the Poisson distribution applies. The Poisson assumption implies that the number of scoring events in non-overlapping time intervals are independent—a reasonable approximation for most sports, though it ignores game-state effects (e.g., a team that is losing may play more aggressively, increasing the scoring rate).

The scoring rates can be decomposed as:

$$\lambda_A = \alpha_A \cdot \delta_B \cdot \gamma, \quad \lambda_B = \alpha_B \cdot \delta_A \cdot \gamma, \quad (21)$$

where  $\alpha_t$  is team  $t$ 's *attack strength* (offensive ability),  $\delta_t$  is team  $t$ 's *defense weakness* (the inverse of defensive ability), and  $\gamma$  is the league-average scoring rate. This decomposition, introduced by Dixon and Coles [18], allows the model to capture the interaction between one team's offense and the other's defense.

Under the Poisson model, the probability that team  $A$  scores exactly  $s_A$  points and team  $B$  scores exactly  $s_B$  points is:

$$P(S_A = s_A, S_B = s_B) = \frac{e^{-\lambda_A} \lambda_A^{s_A}}{s_A!} \cdot \frac{e^{-\lambda_B} \lambda_B^{s_B}}{s_B!}. \quad (22)$$

The win, draw, and loss probabilities are obtained by summing over all score combinations:

$$P(\text{A wins}) = \sum_{s_A > s_B} P(S_A = s_A, S_B = s_B), \quad (23)$$

$$P(\text{draw}) = \sum_{s_A = s_B} P(S_A = s_A, S_B = s_B), \quad (24)$$

$$P(\text{B wins}) = \sum_{s_B > s_A} P(S_A = s_A, S_B = s_B). \quad (25)$$

Under the Poisson model, the total score  $S_A + S_B$  follows a Poisson distribution with rate  $\lambda_A + \lambda_B$  (by the additivity property of independent Poisson variables). The probability that the total exceeds a threshold  $\tau$  (the Over/Under line) is:

$$P(\text{Over}) = 1 - \sum_{k=0}^{\lfloor \tau \rfloor} \frac{e^{-(\lambda_A + \lambda_B)} (\lambda_A + \lambda_B)^k}{k!}. \quad (26)$$



Figure 8 demonstrates the Poisson scoring model with three panels: individual score distributions, total score for Over/Under pricing, and the win probability matrix.




Figure 8: Poisson scoring model. (a) Individual score distributions for two teams with  $\lambda_A = 112$  and  $\lambda_B = 105$ . (b) Total score distribution with Over/Under line at 220.5; the shaded region shows  $P(\text{Over})$ . (c) Win probability heatmap as a function of both teams' expected scores.

Panel (a) shows the individual score distributions. The overlap region between the two distributions represents games where the outcome is uncertain; the non-overlapping tails represent likely blowouts. Panel (b) demonstrates Over/Under pricing: given the O/U line at 220.5, the model computes  $P(\text{Over})$  by summing the probability mass above the line. If  $P(\text{Over})$  differs from the market's implied probability, a +EV bet exists. Panel (c) shows the complete win probability matrix—a core tool used by quantitative sports researchers to identify mispricings across the full space of possible matchups.

### 6.13.2 Application to Point Spread Markets

The Poisson model also enables point spread pricing. The probability that team  $A$  wins by exactly  $m$  points is:

$$P(\text{margin} = m) = \sum_{s_A - s_B = m} P(S_A = s_A, S_B = s_B) = \sum_{k=\max(0, -m)}^{\infty} \frac{e^{-\lambda_A} \lambda_A^{k+m}}{(k+m)!} \cdot \frac{e^{-\lambda_B} \lambda_B^k}{k!}. \quad (27)$$

The probability that team  $A$  covers a spread of  $-S$  (i.e., wins by more than  $S$  points) is:

$$P(\text{cover}) = \sum_{m > S} P(\text{margin} = m). \quad (28)$$

This is directly computable from the Poisson model and provides a second angle for identifying value bets in the point spread market. Professional quant sports researchers routinely compute (28) for every game on the board, comparing the model's cover probability to the market-implied cover probability to identify +EV bets.

## 6.14 Closing Line Value (CLV) Analysis

Closing Line Value (CLV) is widely regarded by professional sports bettors as the single most important metric for evaluating a bettor's long-term profitability. CLV measures whether a bettor consistently gets better odds than the closing line—the final odds available before game time.

**Definition 6.1** (Closing Line Value). *The closing line value for a bet placed at probability  $p_{bet}$  that closes at probability  $p_{close}$  is:*

$$CLV = p_{model} - p_{close}. \quad (29)$$

*A positive CLV indicates the bettor “beat the close”—their prediction was closer to the true probability than the market’s closing assessment.*

The reason CLV is considered the gold standard metric (rather than, say, win rate or ROI) is that CLV is less subject to variance than outcomes-based metrics. A bettor with positive CLV will be profitable in the long run even if their short-term results are negative, because the law of large numbers guarantees that the expected return (which is proportional to CLV) will dominate the variance over a sufficient sample. Conversely, a bettor with negative CLV but positive short-term results is experiencing luck that will eventually regress.

Figure 9 shows the CLV analysis for our Elo model.




Figure 9: Closing Line Value analysis. (a) Scatter plot of model predictions vs. closing line probabilities; points near the diagonal indicate the model agrees with the close. (b) Distribution of CLV values; the mean CLV near zero confirms the model has no systematic advantage over the closing line.

Panel (a) shows that the model’s predictions cluster tightly around the closing line probabilities, with no systematic deviation. Panel (b) confirms that the mean CLV is approximately zero, indicating no systematic edge. This is precisely the result expected under market efficiency: the closing line has already incorporated the information in the Elo model, so the model cannot systematically predict where the closing line will move.

CLV analysis is the primary evaluation metric used by quantitative sports trading firms to assess trader performance. A trader with consistently positive CLV is demonstrating genuine skill, not luck—the sports betting equivalent of consistently generating alpha after transaction costs in financial markets.

### 6.15 Information Edge Decay and Speed-to-Market

A critical concept in professional sports trading—and one that directly parallels high-frequency trading in financial markets—is the *decay of information edge over time*. As new information

becomes available (injury reports, lineup confirmations, weather updates), the market incorporates it into the odds, causing the edge available to an informed bettor to decay toward zero.

Figure 10 illustrates this concept.




Figure 10: Information edge decay timeline. Public data provides a small, persistent edge that decays gradually. Injury news creates a spike of edge at the time of reporting that decays within hours. Lineup confirmations create a sharp, short-lived edge that decays within minutes. The total available edge is the stacked sum of all information sources.

The figure illustrates three information sources:

**Public data edge (navy):** Information contained in team records, historical performance, and Elo ratings. This edge is small and decays slowly because the market has already incorporated most of this information into the opening line. This is the only information source available to our Elo model, explaining why the backtest shows negative returns.

**Injury news edge (coral):** When a key player is reported as injured or ruled out, the line should shift substantially. A trader who can process this information faster than the market can capture the edge before the line moves. The spike at “injury reported” (6 hours before game time in this illustration) shows the sudden availability of edge that decays over the next few hours as the market adjusts.

**Lineup confirmation edge (green):** Starting lineups are typically confirmed 30–60 minutes before game time. A model that can immediately compute the impact of specific lineup combinations can exploit the brief window before the market fully adjusts. This edge decays within minutes.

At professional sports trading operations, traders operate in the final minutes before game time, processing lineup confirmations and making rapid decisions. Researchers build the models that translate lineup information into probability updates. Developers build the low-latency infrastructure that enables execution within seconds. The information decay curve is structurally identical to the alpha decay observed in high-frequency trading: the edge from a new piece of information diminishes as other market participants incorporate it.

## 6.16 Strategy Parameter Sensitivity Analysis

A rigorous quantitative analysis requires understanding how results vary with the strategy parameters. Figure 11 presents a systematic sensitivity analysis.




Figure 11: ROI sensitivity heatmap across edge threshold (horizontal) and Kelly fraction (vertical). Each cell shows the approximate ROI (%) for that parameter combination. Red cells indicate negative ROI; green cells indicate positive ROI. The parameter space is almost entirely red, confirming that the negative result is robust to parameter choice.

The heatmap reveals several important findings:

**No parameter combination achieves consistently positive ROI.** The entire heatmap is predominantly red (negative ROI), confirming that the negative backtest result is not an artifact of suboptimal parameter selection. This is strong evidence for market efficiency: there is no “magic” combination of edge threshold and Kelly fraction that transforms a fundamentally edge-less strategy into a profitable one.

**Higher Kelly fractions amplify losses.** Moving down the heatmap (increasing  $\lambda$  from 10% to 50%) deepens the negative ROI. The reason is that higher Kelly fractions allocate more capital to each bet, and when the expected value is negative, larger bets produce larger losses.

**Higher edge thresholds have ambiguous effects.** Moving right (increasing  $\theta$  from 1% to 10%) reduces the number of bets placed but does not systematically improve quality, because the model’s edges are noise rather than signal (as demonstrated in Figure 4, panel d).

**Implications for researchers:** This type of systematic parameter sweep is a standard component of quantitative research at trading firms. Presenting only the results for one parameter set, without exploring the sensitivity, would be considered incomplete. The heatmap format enables

rapid visual identification of the “sweet spot” (if one exists) and the “danger zones” (parameter combinations that lead to catastrophic losses).

### 6.17 Risk of Ruin Analysis

Risk of ruin—the probability that the bankroll reaches zero—is a critical concern for professional bettors and is directly analogous to the survival analysis conducted by risk management teams at financial institutions.

Figure 12 presents two complementary views of risk.




Figure 12: Risk analysis. (a) Approximate risk of ruin as a function of true edge for different Kelly fractions. Full Kelly has non-trivial ruin probability even at 5% edge; quarter Kelly is much safer. (b) Distribution of maximum drawdowns across 1,000 simulated paths with 53% win probability and quarter-Kelly sizing. The mean maximum drawdown is substantial even with positive edge.

Panel (a) shows that the risk of ruin depends critically on both the true edge and the Kelly fraction. At full Kelly, even a bettor with a genuine 5% edge faces meaningful ruin probability over a season. At quarter Kelly (our backtest setting), the ruin probability drops substantially, but the strategy requires patience—the reduced bet sizes mean slower bankroll growth and longer periods before the edge manifests statistically.

Panel (b) shows the distribution of maximum drawdowns for a *profitable* strategy (53% win rate at  $-110$  odds, giving a positive edge of approximately 1%). Even with positive edge, the mean maximum drawdown is approximately 15–20%, and the 95th percentile exceeds 30%. This means a professional bettor with a genuine 1% edge should expect to lose 30% of their bankroll at some point during the season in 1 out of 20 seasons. Risk management is therefore essential even for profitable strategies—exactly as in financial portfolio management.

Understanding drawdown distributions is essential for position sizing, bankroll management, and psychological resilience. At quantitative trading firms, risk limits are set based on maximum acceptable drawdown, not just expected return. A trader who cannot withstand the inevitable drawdowns will be forced to reduce position sizes (or be forced out) at precisely the wrong time, locking in losses.

## 6.18 Cross-Asset Risk-Adjusted Return Comparison

Figure 13 places sports betting returns in the broader context of financial market returns, using the Sharpe-like ratio as a common metric.




Figure 13: Risk-adjusted return comparison across asset classes and strategies. The S&P 500 (long-term average), hedge fund index, professional sports betting, our Elo model, risk-free rate, and casino house edge are compared. Professional sports bettors achieve Sharpe ratios comparable to top hedge funds.

The figure reveals the following hierarchy of risk-adjusted returns:

**Casino house edge (Sharpe  $\approx 2.5$ ):** The casino has the highest risk-adjusted return because its edge is deterministic (fixed house edge) and diversified across millions of bets per year. The reason casinos are so profitable is not that their edge is large (2–5% per bet) but that it is *consistent* and applied at massive scale.

**Professional sports betting (Sharpe  $\approx 1.5$ ):** Top sports bettors achieve Sharpe ratios comparable to elite hedge funds, reflecting the combination of genuine edge (1–3% per bet) and high bet frequency (hundreds of bets per month). The reason professional sports bettors can achieve such high Sharpe ratios is the binary nature of outcomes: each bet resolves quickly and independently, providing rapid compounding of edge.

**Hedge fund index (Sharpe  $\approx 0.8$ ):** The average hedge fund achieves a Sharpe ratio below 1.0, reflecting the difficulty of generating consistent alpha in highly competitive financial markets.

**S&P 500 (Sharpe  $\approx 0.5$ ):** The equity risk premium provides a Sharpe ratio of approximately 0.5 over the long term, compensating passive investors for bearing systematic risk.

**Our Elo model (Sharpe  $\approx -0.6$ ):** The negative Sharpe ratio confirms that the strategy destroys value on a risk-adjusted basis, losing more per unit of risk than a risk-free investment.

The comparison demonstrates that professional sports betting, executed with quantitative rigor, offers risk-adjusted returns competitive with the best hedge funds. This provides quantitative

context for why major trading firms invest heavily in sports trading infrastructure alongside their traditional financial market operations.

## 6.19 Arbitrage Detection Framework

**Definition 6.2** (Sports Betting Arbitrage). *A guaranteed arbitrage exists when the sum of the best available implied probabilities across different sportsbooks for all outcomes of an event is less than 1:*

$$\sum_{i=1}^k \frac{1}{D_i^{\text{best}}} < 1, \quad (30)$$

where  $D_i^{\text{best}}$  is the best available decimal odds for outcome  $i$  across all sportsbooks, and the sum is over all  $k$  possible outcomes.

When condition (30) is satisfied, a bettor can allocate stakes across outcomes to guarantee a profit regardless of which outcome occurs. The optimal allocation is:

$$s_i = \frac{B}{D_i^{\text{best}} \cdot \sum_{j=1}^k \frac{1}{D_j^{\text{best}}}}, \quad (31)$$

where  $B$  is the total amount to be wagered. The guaranteed profit is:

$$\pi = B \left( \frac{1}{\sum_{j=1}^k \frac{1}{D_j^{\text{best}}}} - 1 \right) > 0. \quad (32)$$

**Numerical example.** FanDuel offers Team A at +150 ( $D = 2.50$ ) while DraftKings offers Team B at −140 ( $D = 1.714$ ). The arbitrage check:  $1/2.50 + 1/1.714 = 0.400 + 0.583 = 0.983 < 1$ . The margin is  $1 - 0.983 = 1.7\%$ , guaranteeing a risk-free profit of 1.7% on the total capital deployed. The optimal stakes (per \$1,000 total):  $s_A = 1000 \times 0.400/0.983 = \$407$  on Team A at FanDuel,  $s_B = 1000 \times 0.583/0.983 = \$593$  on Team B at DraftKings. If Team A wins:  $407 \times 2.50 = \$1017.50$ . If Team B wins:  $593 \times 1.714 = \$1016.40$ . Both outcomes produce approximately \$17 profit on \$1,000 deployed.

The reason arbitrage opportunities are brief (typically lasting minutes) is that automated scanning systems operated by sharp bettors and trading firms detect and exploit them rapidly, moving the lines back into alignment. Building the infrastructure for real-time arbitrage detection—multi-book API integration, sub-second pricing engine, automated execution—is a core competency of sports trading software developers.

## 6.20 Practical Applications: Production System Architecture

This subsection maps the project’s analytical components to the functional roles within a quantitative sports trading operation, illustrating how each module would fit into a production trading system.

### 6.20.1 Research Function

The quantitative researcher is responsible for identifying new predictive features, testing modeling hypotheses, and validating model performance. The following components map to core research activities:

Table 4: Mapping of analytical components to the research function.

| Activity             | Project Component         | Method                                               |
|----------------------|---------------------------|------------------------------------------------------|
| Statistical modeling | Elo system (Sec. 3.3)     | Logistic paired-comparison model with MOV adjustment |
| Hypothesis testing   | Backtest (Sec. 6.4)       | Walk-forward test of +EV hypothesis; null rejected   |
| Model calibration    | Brier/log loss (Sec. 6.5) | Proper scoring rules evaluation                      |
| Feature engineering  | MOV adjustment (10)       | Log-transformed margin as predictive feature         |
| Simulation           | Monte Carlo (Sec. 6.3)    | 10K-path season projections                          |

### 6.20.2 Trading Function

The sports trader makes real-time decisions based on model outputs and market conditions. The following components map to core trading activities:

Table 5: Mapping of analytical components to the trading function.

| Activity            | Project Component        | Method                                          |
|---------------------|--------------------------|-------------------------------------------------|
| Odds interpretation | OddsConverter (Sec. 3.1) | American/decimal/implied probability conversion |
| EV calculation      | EVCalculator (Sec. 3.2)  | Real-time edge identification                   |
| Position sizing     | Kelly (Sec. 3.6)         | Full derivation and fractional Kelly comparison |
| Risk management     | MDD analysis (Sec. 6.4)  | Drawdown tracking and risk quantification       |
| Market structure    | Vig analysis (Sec. 6.7)  | Understanding of market maker economics         |

### 6.20.3 Technology Function

The quantitative developer builds the infrastructure that powers trading operations. The following components map to core engineering activities:



Table 6: Mapping of analytical components to the technology function.

| Activity        | Project Component     | Method                                                      |
|-----------------|-----------------------|-------------------------------------------------------------|
| System design   | 10-class architecture | Modular, single-responsibility classes                      |
| Data pipelines  | DataGenerator         | Synthetic data generation with configurable parameters      |
| Performance     | Complexity analysis   | $O(NG)$ Monte Carlo, $O(1)$ Elo updates                     |
| Visualization   | 18-figure engine      | Publication-quality multi-panel figures                     |
| Reproducibility | Seed management       | Deterministic results via <code>np.random.seed(2026)</code> |

## 6.21 Market Microstructure: Price Discovery in Sports Betting

The mechanics of how sportsbook lines move from the opening odds to the closing line constitute a price discovery process that is structurally identical to price formation in financial markets. Understanding this process is essential for building systems that can identify and exploit transient mispricings.

### 6.21.1 The Price Discovery Process

When a sportsbook opens a line (typically 24–48 hours before game time), the initial odds reflect the book’s internal model. Over the following hours, the line adjusts in response to two forces:

**Sharp money.** Professional bettors with sophisticated models place bets on lines they consider mispriced. These bets carry disproportionate information content because the bettor has already identified a perceived edge. Sportsbooks actively monitor which accounts are “sharp” (consistently profitable) and adjust lines more aggressively in response to their bets. This is directly analogous to *informed order flow* in equity markets, where market makers widen spreads or adjust quotes when they detect informed trading.

**Public money.** Recreational bettors place bets based on team loyalty, media narratives, and gut feelings rather than quantitative models. While each individual public bet carries little information, the aggregate volume creates *inventory risk* for the sportsbook: if 80% of the money is on one side, the book faces a large potential loss. The book adjusts the line to “balance the book”—attracting money on the other side. This is analogous to a market maker adjusting quotes to manage inventory.

Figure 14 illustrates the price discovery process.




Figure 14: Price discovery in a sports betting market. (a) The implied probability (market line) moves from the opening value toward the true probability over 24 hours, driven by sharp money and modulated by public betting noise. The green-shaded area represents the available edge for an informed bettor. (b) Bet volume spikes at market open and near game time; line movements are largest when volume is highest, consistent with Kyle’s lambda model of price impact.

The green-shaded area in panel (a) represents the edge available to an informed bettor who knows the true probability. This area shrinks over time as the line converges toward truth—the

sports betting equivalent of the *efficient market hypothesis in progress*. The key trading insight is that edge is maximized immediately after the line opens (when the market has processed the least information) and decays toward zero as the close approaches.

Panel (b) shows the relationship between volume and line movement. The correlation between these two quantities is a manifestation of *Kyle's lambda*—the price impact coefficient from Albert Kyle's 1985 model of informed trading [19]. In the Kyle model, the market maker cannot distinguish between informed and uninformed order flow, so all orders move the price by an amount proportional to volume. The same mechanism operates in sports betting: the book moves the line in response to volume, regardless of whether the volume is from sharp or public bettors.

### 6.21.2 Closing Line as the Efficient Price

The *closing line value* (CLV) analysis in Section 6.9 is grounded in the empirical finding that the closing line is the most efficient price in the sports betting market—the most accurate publicly available estimate of the true probability. This is analogous to the closing price in equity markets being the most efficient estimate of fair value for a given trading day.

The reason the closing line is efficient is an information aggregation argument: over the 24–48 hours between opening and closing, the line has incorporated information from all market participants—sharp bettors, public bettors, injury reports, weather updates, and the book's own models. Any systematic mispricing in the opening line has been arbitrated away by the time the game starts. A bettor who consistently gets better odds than the closing line (positive CLV) is, by definition, exploiting information that the market has not yet incorporated at the time of the bet.

## 6.22 Bayesian In-Game Updating: Live Betting Models

Live (in-play) betting is the fastest-growing segment of the sports betting market, and the mathematical framework for real-time probability updating is a direct application of Bayesian inference.

### 6.22.1 Mathematical Framework

Let  $\theta$  represent the latent “true” strength difference between two teams, and let  $D_t = \{s_1, s_2, \dots, s_t\}$  denote the observed scoring events through time  $t$ . The posterior win probability after observing data  $D_t$  is:

$$P(\text{A wins} \mid D_t) = \int P(\text{A wins} \mid \theta, D_t) \cdot p(\theta \mid D_t) d\theta, \quad (33)$$

where  $p(\theta \mid D_t)$  is the posterior distribution over strength, updated from the prior  $p(\theta \mid D_0)$  using Bayes' rule:

$$p(\theta \mid D_t) \propto p(D_t \mid \theta) \cdot p(\theta \mid D_0). \quad (34)$$

In practice, the integral (33) is approximated by conditioning on the current score margin  $m_t = s_A(t) - s_B(t)$  and the remaining time  $\tau$ :

$$P(\text{A wins} \mid m_t, \tau) \approx \Phi\left(\frac{m_t}{\sigma\sqrt{\tau}}\right), \quad (35)$$

where  $\Phi$  is the standard normal CDF and  $\sigma$  is the per-unit-time scoring volatility. This approximation treats the scoring process as a Brownian motion—a simplification introduced by Stern [19] for football and subsequently applied to basketball, baseball, and other sports.

The reason the Brownian motion approximation works reasonably well for basketball is that the game consists of many approximately independent scoring events (each possession), and the CLT

ensures that the sum of many small increments converges to a Gaussian process. The approximation breaks down near the end of the game (when teams adjust strategy based on the score), during timeout clusters, and when key players foul out.

Figure 15 shows a simulated basketball game with real-time Bayesian win probability updating.




Figure 15: Bayesian in-game updating for a simulated basketball game. (a) Live win probability for Team A, starting at the pre-game estimate of 62% and fluctuating as the score evolves. The probability converges toward 0 or 1 as the game nears completion. (b) Score margin evolution, with green indicating Team A leading and red indicating Team B leading.

Panel (a) illustrates several important phenomena. Early in the game (first quarter), the win probability fluctuates modestly because each scoring event has limited impact relative to the remaining time. As the game progresses, each scoring event has a larger impact on the win probability because there is less time for the trailing team to recover. Near the end of the game, the win probability approaches 0 or 1 rapidly—the “volatility squeeze” that makes late-game live betting both exciting and mathematically challenging.

Panel (b) shows the score margin, which is the primary input to the Bayesian model. The correlation between the margin evolution (panel b) and the win probability (panel a) is direct but nonlinear: a 5-point lead in the first quarter corresponds to approximately 58% win probability, while the same 5-point lead with 2 minutes remaining corresponds to approximately 90%.

### 6.22.2 Trading Implications

Live betting markets update odds every few seconds, creating a high-frequency trading environment within sports. The key challenges are:

1. **Latency:** The bettor who can compute the posterior probability (35) fastest has the first opportunity to exploit a stale line. At professional trading firms, the model must update within milliseconds of a scoring event being registered by the data feed.
2. **Model accuracy:** The Brownian approximation (35) is a first-order model. More sophisticated approaches incorporate possession-level models, free throw shooting distributions, clock management strategies, and player-specific impact ratings.

3. **Market impact:** Large bets on live markets move the line rapidly. A trader must balance the speed of execution against the risk of moving the price unfavorably before the full position is filled—the same fill-rate vs. market-impact tradeoff faced by algorithmic traders in equity markets.

## 6.23 Portfolio Theory Applied to Sports Betting

A sophisticated bettor does not evaluate each bet in isolation; instead, the entire *book of bets* is treated as a portfolio, and the tools of modern portfolio theory (Markowitz, 1952) are applied to optimize the risk-return profile.

### 6.23.1 Correlation Structure of Game Outcomes

While individual game outcomes are *approximately* independent, several factors introduce correlation:

- **Scheduling:** Teams playing back-to-back games may underperform due to fatigue, creating negative correlation between a team's outcomes on consecutive days.
- **Common opponents:** If Team A beats Team C convincingly, this may update our estimate of Team C's strength, indirectly affecting the prediction for Team C's next game.
- **Weather and venue:** Outdoor sports (NFL, MLB) share weather exposure; multiple games at the same venue (e.g., doubleheaders) share conditions.
- **League-wide trends:** Rule changes, referee assignments, or league-wide shooting trends can affect multiple games simultaneously.

Figure 16 develops the portfolio framework.




Figure 16: Portfolio theory applied to sports betting. (a) Correlation matrix of outcomes for 8 simultaneous NBA games, showing mostly near-zero correlations with a few structurally-induced exceptions. (b) Efficient frontier of bet portfolios, with the optimal Sharpe ratio portfolio marked. (c) Diversification benefit: portfolio risk decreases as  $1/\sqrt{n}$  for independent bets but hits a floor for correlated bets.

### 6.23.2 Optimal Bet Allocation

Given  $n$  available bets with expected return vector  $\boldsymbol{\mu} \in \mathbb{R}^n$  and covariance matrix  $\Sigma \in \mathbb{R}^{n \times n}$ , the Markowitz mean-variance optimization problem is:

$$\max_{\mathbf{w}} \quad \mathbf{w}^\top \boldsymbol{\mu} - \frac{\gamma}{2} \mathbf{w}^\top \Sigma \mathbf{w} \quad \text{s.t.} \quad \mathbf{w} \geq 0, \quad \mathbf{1}^\top \mathbf{w} \leq B, \quad (36)$$

where  $\mathbf{w}$  is the vector of wager amounts,  $\gamma$  is the risk aversion parameter, and  $B$  is the total bankroll. The constraint  $\mathbf{w} \geq 0$  reflects the fact that most bettors cannot “short” a bet (bet against an outcome on the same platform), though cross-book hedging provides an approximate short mechanism.

The efficient frontier in panel (b) of Figure 16 shows the set of portfolios that achieve the maximum expected return for each level of risk. The optimal portfolio (marked with a star) achieves the highest Sharpe ratio. The key insight is that even with modest expected returns per bet (1–5%), the Sharpe ratio of the optimal portfolio can be substantially higher than any individual bet’s Sharpe ratio, due to the diversification benefit of combining approximately independent bets.

Panel (c) quantifies the diversification benefit. For  $n$  independent bets with identical variance  $\sigma^2$ , the portfolio variance is  $\sigma^2/n$ , so the risk decreases as  $1/\sqrt{n}$ . For correlated bets with average pairwise correlation  $\rho$ , the portfolio variance converges to  $\rho\sigma^2$  as  $n \rightarrow \infty$ —a positive floor that cannot be diversified away. This systematic risk component is the sports-betting analogue of *market risk* in financial portfolios.

## 6.24 Signal Analysis: Feature Importance and Half-Life Estimation

A core activity in quantitative research—at sports trading firms and financial institutions alike—is identifying which predictive features carry genuine signal and estimating how quickly that signal decays.

### 6.24.1 Signal Half-Life

**Definition 6.3** (Signal Half-Life). *The half-life of a predictive signal is the time (measured in games, hours, or other units) after which the signal’s predictive power decays to 50% of its initial value:*

$$S(t) = S_0 \cdot 2^{-t/t_{1/2}}, \quad (37)$$

where  $S_0$  is the initial signal strength and  $t_{1/2}$  is the half-life.

Different features have dramatically different half-lives, reflecting how quickly the information they capture is incorporated into the market:

- **Elo rating** ( $t_{1/2} \approx 80$  games): Elo captures long-term team strength, which changes slowly. The signal persists for months but provides little edge because the market has also incorporated historical performance.
- **Recent form** ( $t_{1/2} \approx 25$  games): A team’s performance over the last 10 games captures momentum, injuries, and lineup changes not yet reflected in Elo. The signal decays faster because recent form is inherently transient.
- **Injury news** ( $t_{1/2} \approx 5$  games): When a key player is ruled out, the impact on win probability is immediate and large, but it is also quickly incorporated into the line. The half-life of 5 games reflects the time it takes for all sportsbooks to adjust.

- **Lineup data** ( $t_{1/2} \approx 1$  game): Starting lineups are confirmed 30–60 minutes before tip-off. A model that processes this faster than the market has a brief but potentially large edge.
- **Live game state** ( $t_{1/2} \approx 0.5$  games): The current score and time remaining are the most powerful predictive features during a game, but they are also immediately available to all market participants, giving a half-life measured in seconds.

Figure 17 visualizes the signal landscape.




Figure 17: Signal analysis. (a) Signal half-life curves for six feature types, showing exponential decay at rates that vary by two orders of magnitude. Features with short half-lives (lineup data, live game state) offer the highest potential edge but require the fastest infrastructure to exploit. (b) Feature importance ranked by mutual information with the game outcome, showing that injury-adjusted ratings and Elo rating differences are the strongest static predictors.

Panel (b) ranks features by mutual information (MI) with the game outcome—a non-parametric measure of statistical dependence that captures both linear and nonlinear relationships:

$$I(X; Y) = \sum_{x,y} p(x, y) \log \frac{p(x, y)}{p(x)p(y)}. \quad (38)$$

The ranking reveals that *injury-adjusted ratings* provide the most information about game outcomes (MI = 0.048 bits), followed by *season win percentage* (0.038), *opponent Elo* (0.035), and *Elo rating difference* (0.042). The lowest-information features are *travel distance* (0.008) and *rest days difference* (0.012). The key research question is whether any of these features provides information *not already captured by the closing line*—which is the question of market efficiency addressed throughout this report.

## 6.25 Market Maker Simulation: The Other Side of the Trade

To fully understand sports betting markets, it is essential to model the sportsbook’s perspective. The sportsbook is a market maker: it provides liquidity by offering odds on both sides of every game, earning the vig as compensation for the risk of imbalanced action.

### 6.25.1 Market Maker P&L Model

For a single game with handle (total amount wagered) split as  $H_A$  on Team A and  $H_B$  on Team B, the sportsbook's P&L is:

$$\text{P\&L} = \begin{cases} H_B - H_A \cdot (D_A - 1) & \text{if A wins,} \\ H_A - H_B \cdot (D_B - 1) & \text{if B wins,} \end{cases} \quad (39)$$

where  $D_A$  and  $D_B$  are the decimal odds offered on each side.

If the book is perfectly balanced ( $H_A \cdot D_A = H_B \cdot D_B$ ), the sportsbook earns a guaranteed profit equal to  $\text{vig} \cdot (H_A + H_B)$  regardless of the outcome. In practice, perfect balance is rarely achieved, and the book bears directional risk on every game. The book's expected P&L across many games is positive (due to the vig) but individual games can produce losses.

Figure 18 simulates a sportsbook's season.




Figure 18: Market maker simulation across 500 games. (a) Cumulative P&L showing the steady upward trend produced by the vig, with intermittent drawdowns when the book is caught on the wrong side. (b) Per-game hold percentage distribution, centered near the expected vig of 4.5% but with substantial game-to-game variance.

Panel (a) shows the sportsbook's cumulative P&L. The overall trend is positive—the vig guarantees long-run profitability—but the path is not monotonic. The drawdowns occur when the book has imbalanced action and the outcome favors the heavily-bet side. The maximum drawdown of approximately \$0.5M on a handle of \$50M–\$500M per game illustrates the *inventory risk* that sportsbooks must manage.

Panel (b) shows the per-game hold percentage—the fraction of total handle retained by the book. The mean hold of approximately 4–5% matches the vig, but the variance is substantial: individual games can produce holds of  $-20\%$  (the book loses 20% of the handle) or  $+15\%$  (the book captures far more than the vig). The reason for this variance is the *imbalance*: when public money overwhelms one side, the book is effectively taking a directional bet against the public.

### 6.25.2 Implications for Line-Setting

The market maker simulation reveals why sportsbooks adjust lines: the primary objective is not to predict the outcome correctly but to *minimize inventory risk* by balancing the action. When the



book detects one-sided action (e.g., 80% of the handle on the favorite), it moves the line to make the underdog more attractive, sacrificing theoretical accuracy for risk reduction.

This creates an exploitable dynamic for sharp bettors: when public money pushes a line away from its fair value, sharp bettors can bet against the public and capture the edge. This “fading the public” strategy is one of the oldest and most well-documented edges in sports betting, and it is closely related to the contrarian strategies used in financial markets (buying when others are selling, and vice versa).

Understanding both sides of the market—the bettor’s search for +EV and the book’s risk management—is essential for anyone working in quantitative sports trading, whether on the buy side (quantitative betting operations) or the sell side (sportsbook platforms).

## 7 Conclusion

This project implemented and tested a comprehensive quantitative sports betting analysis system, translating mathematical concepts into working Python code with rigorous backtesting and empirical validation. The principal conclusions are as follows.

**First, Elo models produce well-calibrated probabilities.** The model’s Brier score (0.2312) and log loss (0.6550) demonstrate genuine predictive value for NBA game outcomes, beating the trivial baselines of 0.2500 and 0.6931, respectively. The calibration plot confirms that predicted probabilities match observed win rates across the full range of predictions. The logistic model underlying Elo provides an honest mapping from rating differences to win probabilities, and the margin-of-victory adjustment improves the signal-to-noise ratio of the estimates.

**Second, parlay risk compounds super-linearly.** The effective house edge grows from approximately 5% on single bets to over 60% on 10-leg parlays, providing rigorous quantitative confirmation that sportsbooks earn disproportionate margins on multi-leg bets. The mathematical mechanism—multiplicative compounding of per-leg vig—is formalized in Proposition 3.1 and verified empirically in Figure 1 and Table 1.

**Third, Monte Carlo simulation enables robust scenario analysis.** With 10,000 simulations, the system produces playoff probability estimates with standard errors below 0.5%, demonstrating the technique’s power for uncertainty quantification in sports analytics. The bell-shaped win-total distributions confirm the Central Limit Theorem’s applicability to season-aggregate outcomes.

**Fourth, basic Elo does not overcome market efficiency.** This is the most important finding. The negative backtest ROI (−5.81%) and the absence of correlation between computed edge magnitude and win rate (Figure 4, panel d) demonstrate that a fundamental Elo model, while accurate, does not capture information beyond what is already embedded in closing lines. This is the sports-betting analogue of the efficient market hypothesis: beating the market requires genuine informational or analytical advantages, not just accurate probability estimation. The distinction between model accuracy and betting profitability is a critical pedagogical lesson for quantitative finance students—a lesson that is easier to demonstrate in the sports betting context (where outcomes are binary and observable within hours) than in the financial context (where returns are continuous and outcomes may take months or years to materialize).

**Fifth, Kelly Criterion provides rigorous position sizing.** The comparison of Kelly fractions (Figure 6) demonstrates the growth-volatility tradeoff quantitatively. Full Kelly maximizes long-run growth but produces extreme path volatility; fractional Kelly sacrifices growth for stability. The severe consequences of overbetting—visible in the full-Kelly trajectory’s deep drawdowns—justify the industry practice of using fractional Kelly (typically 25–50% of the optimal fraction) in both sports betting and financial portfolio management.

**Sixth, the mathematical toolkit is identical across domains.** The project confirms that the mathematical toolkit for sports betting—probability theory, Bayesian updating (Elo), Monte Carlo simulation, position sizing (Kelly), and risk management (drawdown analysis)—is structurally identical to quantitative finance. The transition from one domain to the other is not a change in methodology but a change in data source and payout structure.

**Future work** could extend this system with: (a) player-level features such as injury-adjusted ratings, rest-day effects, and travel distance; (b) live Bayesian updating for in-game betting, where probabilities shift every possession; (c) integration with real-time odds APIs from platforms like FanDuel and DraftKings for live arbitrage detection; (d) cross-market arbitrage detection by monitoring simultaneous odds across multiple sportsbooks; and (e) deep learning models (e.g., recurrent neural networks) that capture temporal dependencies in team performance. The modular architec-

ture of the current system is designed to support these extensions without requiring a rewrite of the core infrastructure.

## References

- [1] G. W. Brier, “Verification of forecasts expressed in terms of probability,” *Monthly Weather Review*, vol. 78, no. 1, pp. 1–3, 1950.
- [2] A. E. Elo, *The Rating of Chessplayers, Past and Present*. Arco Publishing, 1978.
- [3] J. L. Kelly, “A new interpretation of information rate,” *Bell System Technical Journal*, vol. 35, no. 4, pp. 917–926, 1956.
- [4] N. Silver, “Introducing NFL Elo ratings,” *FiveThirtyEight*, Sept. 2014. [Online]. Available: <https://fivethirtyeight.com/features/introducing-nfl-elo-ratings/>
- [5] E. O. Thorp, *Beat the Dealer: A Winning Strategy for the Game of Twenty-One*. Vintage Books, 1966.
- [6] E. O. Thorp, *A Man for All Markets: From Las Vegas to Wall Street, How I Beat the Dealer and the Market*. Random House, 2017.
- [7] American Gaming Association, “Commercial gaming revenue tracker,” 2025. [Online]. Available: <https://www.americangaming.org/>
- [8] S. A. Kovalchik, “Searching for the GOAT of tennis win prediction,” *Journal of Quantitative Analysis in Sports*, vol. 12, no. 3, pp. 127–138, 2016.
- [9] M. E. Glickman, “Parameter estimation in large dynamic paired comparison experiments,” *Journal of the Royal Statistical Society: Series C*, vol. 48, no. 3, pp. 377–394, 1999.
- [10] H. Rue and O. Salvesen, “Prediction and retrospective analysis of soccer matches in a league,” *The Statistician*, vol. 49, no. 3, pp. 399–418, 2000.
- [11] R. T. Stefani, “The methodology of officially recognized international sports rating systems,” *Journal of Quantitative Analysis in Sports*, vol. 7, no. 4, 2011.
- [12] K. P. Murphy, *Machine Learning: A Probabilistic Perspective*. MIT Press, 2012.
- [13] M. Baxter and A. Rennie, *Financial Calculus: An Introduction to Derivative Pricing*. Cambridge University Press, 1996.
- [14] M. Haugh, “The Monte Carlo framework, examples from finance and generating correlated random variables,” *Columbia University Lecture Notes*, 2004.
- [15] Supreme Court of the United States, “Murphy v. National Collegiate Athletic Association,” 584 U.S. 453, 2018.
- [16] L. M. Hvattum and H. Arntzen, “Using ELO ratings for match result prediction in association football,” *International Journal of Forecasting*, vol. 26, no. 3, pp. 460–470, 2010.
- [17] A. C. Constantinou, N. E. Fenton, and W. Marsh, “Pi-football: A Bayesian network model for forecasting Association Football match outcomes,” *Knowledge-Based Systems*, vol. 36, pp. 322–339, 2012.
- [18] M. J. Dixon and S. G. Coles, “Modelling association football scores and inefficiencies in the football betting market,” *Journal of the Royal Statistical Society: Series C*, vol. 46, no. 2, pp. 265–280, 1997.

- [19] H. S. Stern, “A Brownian motion model for the progress of sports scores,” *Journal of the American Statistical Association*, vol. 89, no. 427, pp. 1128–1134, 1994.

## A Key Implementation Code Excerpts

This appendix provides the core algorithmic implementations from `sports_betting_analysis.py`. The full source code (900+ lines) is provided as a separate deliverable.

### A.1 Odds Converter

```
class OddsConverter:
    @staticmethod
    def american_to_implied(odds: float) -> float:
        if odds < 0:
            return abs(odds) / (abs(odds) + 100.0)
        else:
            return 100.0 / (odds + 100.0)

    @staticmethod
    def american_to_decimal(odds: float) -> float:
        if odds < 0:
            return 1.0 + 100.0 / abs(odds)
        else:
            return 1.0 + odds / 100.0

    @staticmethod
    def compute_vig(odds_a: float, odds_b: float) -> float:
        p_a = OddsConverter.american_to_implied(odds_a)
        p_b = OddsConverter.american_to_implied(odds_b)
        return (p_a + p_b - 1.0) * 100.0
```

### A.2 Elo Rating System

```
class EloSystem:
    def __init__(self, k_factor=20.0, home_advantage=100.0,
                 initial_rating=1500.0, season_reversion=0.75):
        self.K = k_factor
        self.HOME_ADV = home_advantage
        self.INIT = initial_rating
        self.REVERT = season_reversion
        self.ratings = {}

    def expected_score(self, rating_a, rating_b, home_a=False):
        adj_a = rating_a + (self.HOME_ADV if home_a else 0)
        return 1.0 / (1.0 + 10.0 ** ((rating_b - adj_a) / 400.0))

    def update(self, team_a, team_b, score_a,
              home_a=False, mov=1.0):
        r_a = self._get_rating(team_a)
        r_b = self._get_rating(team_b)
        e_a = self.expected_score(r_a, r_b, home_a)
```

```

e_b = 1.0 - e_a
score_b = 1.0 - score_a
k_adj = self.K * np.log(max(mov, 1) + 1)
self.ratings[team_a] = r_a + k_adj * (score_a - e_a)
self.ratings[team_b] = r_b + k_adj * (score_b - e_b)

```

### A.3 Kelly Criterion

```

class KellyCriterion:
    @staticmethod
    def full_kelly(true_prob, american_odds):
        dec = OddsConverter.american_to_decimal(american_odds)
        b = dec - 1.0
        p = true_prob
        q = 1.0 - p
        f = (b * p - q) / b
        return max(f, 0.0)

    @staticmethod
    def fractional_kelly(true_prob, american_odds, fraction=0.5):
        return fraction * KellyCriterion.full_kelly(
            true_prob, american_odds)

    @staticmethod
    def bankroll_simulation(initial, bets, kelly_frac=0.5):
        bankroll = [initial]
        current = initial
        for true_p, odds, outcome in bets:
            f = KellyCriterion.fractional_kelly(
                true_p, odds, kelly_frac)
            wager = current * f
            dec = OddsConverter.american_to_decimal(odds)
            if outcome == 1:
                current += wager * (dec - 1.0)
            else:
                current -= wager
            bankroll.append(max(current, 0))
        return bankroll

```

### A.4 Monte Carlo Simulator

```

class MonteCarloSimulator:
    def __init__(self, elo_system, n_simulations=10000):
        self.elo = elo_system
        self.N = n_simulations

    def simulate_game(self, team_a, team_b, home_a=False):
        p_a = self.elo.win_probability(team_a, team_b, home_a)

```

```

        return team_a if np.random.random() < p_a else team_b

def simulate_season(self, schedule):
    wins = defaultdict(int)
    for home, away in schedule:
        winner = self.simulate_game(home, away, home_a=True)
        wins[winner] += 1
    return dict(wins)

def run(self, schedule):
    all_results = defaultdict(list)
    for _ in range(self.N):
        season = self.simulate_season(schedule)
        for team in set(t for game in schedule for t in game):
            all_results[team].append(season.get(team, 0))
    return dict(all_results)

def playoff_probability(self, results, threshold=42):
    probs = {}
    for team, wins in results.items():
        probs[team] = np.mean([w >= threshold for w in wins])
    return dict(sorted(probs.items(), key=lambda x: -x[1]))

```

## A.5 Expected Value Calculator

```

class EVCalculator:
    @staticmethod
    def expected_value(true_prob, american_odds, stake=100.0):
        decimal = OddsConverter.american_to_decimal(american_odds)
        ev = true_prob * (decimal - 1.0) * stake \
            - (1.0 - true_prob) * stake
        return ev

    @staticmethod
    def edge(true_prob, implied_prob):
        return (true_prob - implied_prob) * 100.0

    @staticmethod
    def breakeven_prob(american_odds):
        return OddsConverter.american_to_implied(american_odds)

```

## A.6 Model Calibration

```

class MarketEfficiency:
    @staticmethod
    def calibration_test(predictions, outcomes, n_bins=10):
        preds = np.array(predictions)
        outs = np.array(outcomes)

```



```
bins = np.linspace(0, 1, n_bins + 1)
results = {}
for i in range(n_bins):
    mask = (preds >= bins[i]) & (preds < bins[i+1])
    if mask.sum() > 0:
        center = (bins[i] + bins[i+1]) / 2.0
        results[center] = (preds[mask].mean(),
                           outs[mask].mean(),
                           int(mask.sum()))

return results

@staticmethod
def brier_score(predictions, outcomes):
    preds = np.array(predictions)
    outs = np.array(outcomes)
    return np.mean((preds - outs) ** 2)

@staticmethod
def log_loss(predictions, outcomes, eps=1e-15):
    preds = np.clip(predictions, eps, 1 - eps)
    preds = np.array(preds)
    outs = np.array(outcomes)
    return -np.mean(outs * np.log(preds)
                    + (1-outs) * np.log(1-preds))
```